

**Directions:** Type or clearly handwrite your solutions to each of the following exercises. Partial credit cannot be given unless all work is shown. You may work in groups provided that each person takes responsibility for understanding and writing out the solutions. Additionally, you must give proper credit to your collaborators by providing their names on the line below (if you worked alone, write “No Collaborators”):

COLLABORATORS:

## Note to Self

### Confidence Intervals

Formula:

$$\text{Confidence Interval} = (\bar{Y}_1 - \bar{Y}_2) \pm t_{(df-2, 1-\frac{\alpha}{2})} S_p \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}$$

Where:

$$S_p = \frac{(n_1 - 1)S_1^2 + (n_2 - 1)S_2^2}{n_1 + n_2 - 2}$$

Calculation

```
# Load Data
library(readr)
smsspeed_1 <- read_csv("C:/Users/samue/OneDrive/Desktop/Iowa_State_PS/STAT 5000/Labs/Lab 3/smsspeed-1.csv")
```

```
## Rows: 30 Columns: 4
## -- Column specification -----
## Delimiter: ","
## chr (1): AgeGroup
## dbl (3): Age, Own Phone, Control
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
data1 <- smsspeed_1 %>%
  filter(smsspeed_1$AgeGroup == "Over30")

data2 <- smsspeed_1 %>%
  filter(smsspeed_1$AgeGroup == "Teens")

sampleMean1 <- mean(data1$`Own Phone`)
sampleMean2 <- mean(data2$`Own Phone`)
difference <- sampleMean1 - sampleMean2
difference
```

```
## [1] 44.012
```

```
#Step 2: Finding standard deviation
s1 <- sd(data1$`Own Phone`)
s2 <- sd(data2$`Own Phone`)

#Step 3: Finding sample size
n1 <- length(data1$`Own Phone`)
n2 <- length(data2$`Own Phone`)

numerator <- (n1-1)*(s1^2) + (n2-1)*(s2^2)
denom <- n1 + n2 - 2
pooled <- sqrt( numerator / denom )
sqrtFactor <- sqrt(1/n1 + 1/n2)

tStatDf <- 2.0484

rightSide <- tStatDf*pooled*sqrtFactor

pooled
```

```
## [1] 18.51069
```

```
rightSide
```

```
## [1] 13.84544
```

```
lb <- difference - rightSide
ub <- difference + rightSide

lb
```

```
## [1] 30.16656
```

```
ub
```

```
## [1] 57.85744
```

## Assignment

1.

+21

: A major medical center in the Northeastern U.S. conducted a study looking at blood cholesterol levels and incidence of heart attack. Below are summary statistics of blood cholesterol levels from 16 people who had a heart attack and 20 people who did not have a heart attack.

Group	Sample Size ( $n$ )	Sample Mean ( $\bar{y}$ )	Sample Std. Dev. ( $s$ )
Heart Attack (1)	$n_1 = 16$	$\bar{y}_1 = 265.4$	$s_1 = 43.645$
No Heart Attack (0)	$n_2 = 20$	$\bar{y}_2 = 193.1$	$s_2 = 21.623$

1. List the assumptions needed to properly use the  $t$ -based confidence interval.

- $Y_{1,1}, Y_{1,2}, \dots, Y_{1,16}$  are i.i.d.  $N(\mu_1, \sigma^2)$
- $Y_{2,1}, Y_{2,2}, \dots, Y_{2,20}$  are i.i.d.  $N(\mu_2, \sigma^2)$
- Population variances between the two groups are equal (both have  $\sigma^2$  variance)
- Independence between observations between groups, i.e.  $Y_{1,i}$  and  $Y_{2,j}$  are independent for all  $i$  and  $j$

2. Using the formula from lecture, compute a 95% confidence interval "by hand" for the population difference in mean cholesterol level between the heart attack and the no heart attack groups.

```
# df <- 16 + 20 - 2
# tStatistic <- 34, 0.975

tStatistic <- function(n1, n2, mu1, mu2, sd1, sd2) {  numerator1 <- (n1-1)*(sd1^2) + (n2-1)*(sd2^2)
  denom1 <- n1 + n2 - 2
  pooled <- sqrt( numerator1 / denom1 )

  numerator2 <- mu1 - mu2
  denom2 <- pooled * (sqrt( (1/n1) + (1/n2) ) )
  tStatOutput <- numerator2 / denom2
  tStatOutput
}

ciEstimate <- function(n1, n2, mu1, mu2, sd1, sd2, tStat) {

  numerator <- (n1-1)*(sd1^2) + (n2-1)*(sd2^2)
  denom <- n1 + n2 - 2

  pooled <- sqrt( numerator / denom )
  sqrtFactor <- sqrt(1/n1 + 1/n2)
  rightSide <- tStat*pooled*sqrtFactor

  difference <- mu1 - mu2

  lb <- difference - rightSide
  ub <- difference + rightSide
}
```

```

output <- c(lb, ub)
output

}

tStatistic(n1 = 16, n2 = 20, mu1 = 265.4, mu2 = 193.1, sd1 = 43.645, sd2 = 21.623)

## [1] 6.494355

ciEstimate(n1 = 16, n2 = 20, mu1 = 265.4, mu2 = 193.1, sd1 = 43.645, sd2 = 21.623, tStat = 6.494355)

## [1] 4.110884e-06 1.446000e+02

```

3. Interpret the confidence interval in the context of the study.

We are 95% confidence the true difference between the two groups (those who had heart attacks compared to those who didn't) is between -28.29 to 116.31, meaning we believe those who had a heart attack on average have between 116.31 higher to 28.29 lower blood pressure compared to those who didn't have a heart attack. Of note, this 95% confidence interval includes zero.

4. Suppose that the researchers want to replicate the study, varying the targeted sample sizes in each group to obtain the *best* (or least variable) estimate of the difference in group means, given the constraint that they can only afford to collect information from 50 total participants. How many subjects should they recruit for each group? Fill in the table below to help you answer the question (as an example, one of the solutions is already provided).

$$Var(\bar{Y}_1 - \bar{Y}_2) = \sigma^2 \left( \frac{1}{n_1} + \frac{1}{n_2} \right)$$

$n_1$	$n_2$	$Var(\bar{Y}_1 - \bar{Y}_2) = \sigma^2 \left( \frac{1}{n_1} + \frac{1}{n_2} \right) \setminus$
1	49	$1.020\sigma^2$
5	45	$0.222\sigma^2$
10	40	$0.125\sigma^2$
20	30	$0.083\sigma^2$
25	25	$0.08\sigma^2$
30	20	$0.083\sigma^2$
40	10	$0.125\sigma^2$
45	5	$0.222\sigma^2$
49	1	$1.020\sigma^2$

They should evenly divide the 50 total participants into two groups, making two groups of 25.

2.

+10

: Refer to the data set `cholesterol.csv` (posted in Canvas). This file contains data on 18 randomly sampled individuals diagnosed with high cholesterol who replaced butter in their diets with a brand (A or B) of margarine. The brand of margarine was randomized. Their doctors recorded their blood cholesterol levels at the beginning of the experiment, after four weeks of their diets, and again after eight weeks. The researchers are interested in exploring the question:

How big is the difference in mean cholesterol reduction after 8 weeks between brand A and B for subjects who replaced butter in their diet with margarine?

1. Compute the 99% confidence interval in R and provide a screenshot of the output.

```
library(readr)
library(dplyr)
cholesterol <- read_csv("cholesterol.csv")

## Rows: 18 Columns: 5
## -- Column specification -----
## Delimiter: ","
## chr (1): Margarine
## dbl (4): ID, Before, After4weeks, After8weeks
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

differenceDf <- cholesterol |>
  mutate(difference = Before - After8weeks)

brandAdf <- differenceDf |>
  filter(Margarine == "A")
brandBdf <- differenceDf |>
  filter(Margarine == "B")

cholTstatistic <- tStatistic(n1 = nrow(brandAdf), n2 = nrow(brandBdf), mu1 = mean(brandAdf$difference),
cholTstatistic

## [1] -2.152855

ciEstimate(n1 = nrow(brandAdf),
  n2 = nrow(brandBdf),
  mu1 = mean(brandAdf$difference),
  mu2 = mean(brandBdf$difference),
  sd1 = sd(brandAdf$difference),
  sd2 = sd(brandBdf$difference),
  tStat = cholTstatistic)

## [1] 0.0000000 -0.3288889
```

2. Interpret the confidence interval in the context of the study.

We are 95% confident that the true difference in blood cholesterol levels between those who consumed Brand A vs. Brand B is between 43.848 and 44.176; this is further support of the finding that consuming Brand A margarine increases blood cholesterol levels compared with consuming Brand B margarine.

3.

+9

: Suppose that the researchers want to replicate the cholesterol study. Help them achieve their study design goals by performing the following sample size determinations in R (provide screenshot of output) or by hand-calculation (show work):

### Standard Error

Formula

Function

### Confidence Intervals

Formula

Function

```
n0Est <- function(){
  numerator <- (n1-1)*(sd1^2) + (n2-1)*(sd2^2)
  denom <- n1 + n2 - 2
  pooled <- sqrt( numerator / denom )

  # z statistic
  top <-
  # w, width
  bottom <-
  output <- 8 * (top / bottom)^2
  output
}
```

```
nEst <- function(){
  numerator <- (n1-1)*(sd1^2) + (n2-1)*(sd2^2)
  denom <- n1 + n2 - 2
  pooled <- sqrt( numerator / denom )

  # t statistic using n0Est
  top <-
  # w, width
  bottom <-
  output <- 8 * (top / bottom)^2
  output
}
```

### Hypothesis Tests

Formula

Function

1. Given an approximate pooled sample standard deviation of  $S_p = 0.16$ , what sample size is needed in each of two equally-sized treatment groups in order for the standard error of the difference in average cholesterol reduction to be no more than 0.02?
2. Given an approximate pooled sample standard deviation of  $S_p = 0.16$ , what sample size is needed in each of two equally-sized treatment groups in order for the width of a 95% confidence interval for the true difference in mean cholesterol reduction to be no more than 0.04?
3. Given an approximate pooled sample standard deviation of  $S_p = 0.16$  and an effect size of  $\delta = 0.03$ , what sample size is needed in each of two equally-sized treatment groups in order for a level  $\alpha = 0.05$  two-sided test to have 80% power?

4.

+10

: Refer to the data set `birthweight.csv` (posted in Canvas) containing the weights of a random sample of babies born at a certain hospital and information about whether the birth mother was a smoker.

1. Using R, find a 90% confidence interval for the difference in average birthweight between babies born to mothers who smoke versus non-smoking mothers, and provide a screenshot of the output.

```
birthweight <- read_csv("birthweight.csv")

## Rows: 42 Columns: 2
## -- Column specification -----
## Delimiter: ","
## dbl (2): Birthweight, smoker
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.

birthweight1 <- birthweight |>
  filter(smoker == 0)
birthweight2 <- birthweight |>
  filter(smoker == 1)

t.test(birthweight)

##
## One Sample t-test
##
## data: birthweight
## t = 11.658, df = 83, p-value < 2.2e-16
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
## 1.591046 2.245620
## sample estimates:
## mean of x
## 1.918333

ciEstimate(n1 = nrow(birthweight1),
           n2 = nrow(birthweight2),
           mu1 = mean(birthweight1$Birthweight),
           mu2 = mean(birthweight2$Birthweight),
           sd1 = sd(birthweight1$Birthweight),
           sd2 = sd(birthweight2$Birthweight),
           tStat = t.test(birthweight)$statistic)

##          t          t
## -1.715174  2.465992
```

2. Interpret the confidence interval in the context of the study.

Total: 50 points # correct: %: