

# PS6

Sam Olson

## Q1

Consider the dataset `pigs` provided in the R package `emmeans`. The data can be accessed in R with the following commands.

```
library(emmeans)
attach(pigs)
```

To learn a more about the data, type `?pigs` at the R prompt. For the purposes of this problem, use the natural logarithm of the variable `conc` as the response. Consider both `source` and `percent` as categorical factors. Assume the cell-means model with one unrestricted treatment mean for each combination of `source` and `percent`.

```
lnConc <- log(pigs$conc)
pigs$percent <- factor(pigs$percent)
class(pigs$percent)
```

```
## [1] "factor"
```

a)

Generate an ANOVA table with Type I (sequential) sums of squares for `source`, `percent`, `source × percent`, `error`, and `corrected total`. In addition to sums of squares, your ANOVA table should include degrees of freedom, mean squares, F statistics, and p-values where appropriate.

Type 1 Sums of Squares are the default given when using the `anova` function in R. Hence:

```
baseDat <- lm(lnConc ~ source + percent + source*percent, data=pigs)
type1Dat <- anova(baseDat)
type1_df <- data.frame(type1Dat)

total <- c(sum(type1_df$`Df`), sum(type1_df$`Sum.Sq`), NA, NA, NA)
type1_df <- rbind(type1_df, total)
rownames(type1_df)[nrow(type1_df)] <- "Corrected Total"

print(type1_df, digits = 6)
```

##		Df	Sum.Sq	Mean.Sq	F.value	Pr..F.
##	source	2	0.6301386	0.3150693	23.311284	1.34314e-05
##	percent	3	0.3173579	0.1057860	7.826871	1.69937e-03
##	source:percent	6	0.0750848	0.0125141	0.925893	5.01099e-01
##	Residuals	17	0.2297676	0.0135157	NA	NA
##	Corrected Total	28	1.2523490	NA	NA	NA

b)

Generate an ANOVA table with Type II sums of squares for `source`, `percent`, `source × percent`, `error`, and `corrected total`. In addition to sums of squares, your ANOVA table should include degrees of freedom, mean squares, F statistics, and p-values where appropriate.

The `car` package allows more flexible usage of Types of Sums of Squares used for anova. Hence:

```
library(car)
type2Dat <- car::Anova(baseDat, type = 2)

type2_df <- as.data.frame(type2Dat)

# Compute Mean Square (MS) = SS / df
type2_df$MeanSq <- type2_df$`Sum Sq` / type2_df$Df

# Reorder columns to match anova() output
type2_df <- type2_df[, c("Df", "Sum Sq", "MeanSq", "F value", "Pr(>F)")]

total <- c(sum(type2_df$`Df`), sum(type2_df$`Sum Sq`), NA, NA, NA)
type2_df <- rbind(type2_df, total)
rownames(type2_df)[nrow(type2_df)] <- "Corrected Total"
# Print with higher precision
print(type2_df, digits = 6)
```

	Df	Sum Sq	MeanSq	F value	Pr(>F)
## source	2	0.7647592	0.3823796	28.291424	3.90135e-06
## percent	3	0.3173579	0.1057860	7.826871	1.69937e-03
## source:percent	6	0.0750848	0.0125141	0.925893	5.01099e-01
## Residuals	17	0.2297676	0.0135157	NA	NA
## Corrected Total	28	1.3869696	NA	NA	NA

c)

Generate an ANOVA table with Type III sums of squares for `source`, `percent`, `source × percent`, `error`, and `corrected total`. In addition to sums of squares, your ANOVA table should include degrees of freedom, mean squares, F statistics, and p-values where appropriate.

Quick note on the continued use of `car::Anova` for this problem: By default R uses the baseline constraint when making these calculations, and in R the baseline is the combination of the first factor levels. So effectively we just need to change the type of contrast being used when calculating the Anova table from treatment aka `contr.treatment` to sums aka `contra.sum` (which corresponds to a sum-to-zero constraint) and proceed, bearing in mind Type 3 Sums of Squares are ‘marginal’ Sums of Squares, hence the reason for this adjustment.

see: <https://rdoodles.rbind.io/2020/10/type-3-anova-in-r-an-easy-way-to-publish-wrong-tables/>

```
library(car)

# Contrast change from treatment to sum
contrasts(pigs$source) <- contr.sum
contrasts(pigs$percent) <- contr.sum
# could have directly added the above in the contrast argument to the lm function
model <- lm(lnConc ~ source * percent, data = pigs)

# Call Anova with updated contrasts
type3Dat <- car::Anova(model, type = 3)

type3_df <- as.data.frame(type3Dat)
type3_df$MeanSq <- type3_df$`Sum Sq` / type3_df$Df
type3_df <- type3_df[, c("Df", "Sum Sq", "MeanSq", "F value", "Pr(>F)")]

# Want to remove Intercept from calculation
total_Df <- sum(type3_df[-1,]$Df, na.rm = TRUE)
total_SS <- sum(type3_df[-1,]$`Sum Sq`, na.rm = TRUE)

total <- c(total_Df, total_SS, NA, NA, NA)
type3_df <- rbind(type3_df, total)
rownames(type3_df)[nrow(type3_df)] <- "Corrected Total"

print(type3_df[-1,], digits = 6)
```

##		Df	Sum Sq	MeanSq	F value	Pr(>F)
##	source	2	0.8178765	0.4089383	30.256440	2.50691e-06
##	percent	3	0.3330420	0.1110140	8.213681	1.34768e-03
##	source:percent	6	0.0750848	0.0125141	0.925893	5.01099e-01
##	Residuals	17	0.2297676	0.0135157	NA	NA
##	Corrected Total	28	1.4557709	NA	NA	NA

d)

Find LSMeans for source and percent.

```
library(tidyverse)

summary_table <- pigs |>
  group_by(source, percent) |>
  mutate(count = n())

print(summary_table, digits = 6)

## # A tibble: 29 x 4
## # Groups:   source, percent [12]
##   source percent  conc count
##   <fct>  <fct>   <dbl> <int>
## 1 fish    9      27.8     2
## 2 fish    9      23.7     2
## 3 fish   12      31.5     3
## 4 fish   12      28.5     3
## 5 fish   12      32.8     3
## 6 fish   15      34      2
## 7 fish   15      28.3     2
## 8 fish   18      30.6     3
## 9 fish   18      32.7     3
## 10 fish  18      33.7     3
## # i 19 more rows
```

Our data is not balanced, so need to explicitly calculate LSmeans (they are not the same as the OLS estimates).

```
library(tidyverse)
library(emmeans)

modDat <- lm(lnConc ~ source * percent, data = pigs)
lsmeans_table <- emmeans(modDat, ~ source * percent)
lsmeans_df <- as.data.frame(lsmeans_table)

lsmeans_wide <- lsmeans_df |>
  select(source, percent, emmean) |>
  pivot_wider(names_from = percent, values_from = emmean)

lsmeans_wide

## # A tibble: 3 x 5
##   source   '9'  '12'  '15'  '18'
##   <fct> <dbl> <dbl> <dbl> <dbl>
## 1 fish   3.25  3.43  3.43  3.48
## 2 soy    3.54  3.68  3.67  3.76
## 3 skim   3.56  3.76  3.90  4.09
```

We then sum across row, column for the source, percent LSmeans respectively.

```
lsmeans_wide <- lsmeans_wide |>
  mutate(Row_Avg = rowMeans(select(lsmeans_wide, -source), na.rm = TRUE))

lsmeans_wide[-c(2:5)]
```

```
## # A tibble: 3 x 2
##   source Row_Avg
##   <fct>   <dbl>
## 1 fish     3.40
## 2 soy      3.66
## 3 skim     3.83
```

```
# Compute column-wise averages (across sources)
col_avg <- lsmeans_wide |>
  summarise(across(-source, mean, na.rm = TRUE)) |>
  mutate(source = "Column_Avg")

col_avg
```

```
## # A tibble: 1 x 6
##   '9' '12' '15' '18' Row_Avg source
##   <dbl> <dbl> <dbl> <dbl>   <dbl> <chr>
## 1  3.45  3.62  3.67  3.78     3.63 Column_Avg
```

In Summary, our LSmeans are:

Percent:

- 9%: 3.448
- 12%: 3.625
- 15%: 3.670
- 18%: 3.775

Source:

- Fish: 3.396
- Soy: 3.662
- Skim: 3.830

e)

Consider simplifying the model so that `percent` is treated like a quantitative variable with linear effects on `log(conc)` and linear interactions; i.e.,

```
lm(y ~ source + percent + source:percent)
```

where `y=log(conc)` and `percent` is numeric. Does such a model fit adequately relative to the cell-means model? Conduct a lack of fit test and report the results.

Back to the land of vanilla `anova`. We can just put the two models in and compare directly.

```
# Need to readjust based on messing with contrasts for Type 3 Sums of Squares
contrasts(pigs$source) <- contr.sum
contrasts(pigs$percent) <- contr.sum
# Reduced model: percent as numeric
reduced <- lm(lnConc ~ source + as.numeric(percent) + source:as.numeric(percent), data = pigs)

# Full model: percent as factor
full <- lm(lnConc ~ source * factor(percent), data = pigs)

anova(reduced, full)

## Analysis of Variance Table
##
## Model 1: lnConc ~ source + as.numeric(percent) + source:as.numeric(percent)
## Model 2: lnConc ~ source * factor(percent)
##   Res.Df    RSS Df Sum of Sq    F Pr(>F)
## 1      23 0.26291
## 2      17 0.22977  6  0.033141 0.4087 0.8631

# order is important for positive df, but statistical tests yield same results (F stat and p-value)
# anova(full, reduced)
```

We have evidence to support the reduced model being adequate in lieu of the more complex, cell means model.

f)

The reduced model fit in part e) implies that, for each **source**, there is a linear relationship between the expected log concentration and percentage. Based on the fit of the reduced model in part e), provide the estimated linear relationship for each **source**.

General Form of Model:

$$y_{ijk} = \mu + \alpha_i + \beta x_{ij} + \gamma_i x_{ij} + \epsilon_{ijk}$$

Where:

$$E[y_{ijk}] = \mu + \alpha_i + \beta x_{ij} + \gamma_i x_{ij}$$

And where we use the following as our estimated effects:

```
library(emmeans)
attach(pigs)
pigs$lnConc <- log(pigs$conc)
pigs$percent <- as.numeric(as.character(pigs$percent))

reducedFish <- lm(lnConc ~ percent,
                  data = subset(x = pigs, source == "fish"))

reducedSoy <- lm(lnConc ~ percent,
                  data = subset(x = pigs, source == "soy"))

reducedSkim <- lm(lnConc ~ percent,
                   data = subset(x = pigs, source == "skim"))

summary(reducedFish)

##
## Call:
## lm(formula = lnConc ~ percent, data = subset(x = pigs, source ==
##      "fish"))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.140840 -0.061335  0.004943  0.065602  0.120815
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   3.11642    0.12083   25.791 5.48e-09 ***
## percent       0.02110    0.00851    2.479  0.0382 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0899 on 8 degrees of freedom
## Multiple R-squared:  0.4345, Adjusted R-squared:  0.3638
## F-statistic: 6.147 on 1 and 8 DF, p-value: 0.03815
```



```
summary(reducedSoy)
```

```
##
## Call:
## lm(formula = lnConc ~ percent, data = subset(x = pigs, source ==
##      "soy"))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.169329 -0.022792 -0.003073  0.050349  0.107387
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.368104    0.111084  30.320 1.52e-09 ***
## percent      0.021748    0.008586   2.533  0.0351 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.0798 on 8 degrees of freedom
## Multiple R-squared:  0.4451, Adjusted R-squared:  0.3757
## F-statistic: 6.417 on 1 and 8 DF, p-value: 0.03509
```

```
summary(reducedSkim)
```

```
##
## Call:
## lm(formula = lnConc ~ percent, data = subset(x = pigs, source ==
##      "skim"))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.195329 -0.123058 -0.001485  0.132921  0.169293
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.04920    0.20579  14.817 1.53e-06 ***
## percent      0.05796    0.01622   3.574  0.00905 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1451 on 7 degrees of freedom
## Multiple R-squared:  0.646, Adjusted R-squared:  0.5954
## F-statistic: 12.77 on 1 and 7 DF, p-value: 0.009051
```

Based on the above output, our estimates are:

**Fish**

$$\hat{y} = 3.11642 + 0.02110x_{1j}$$

**Soy**

$$\hat{y} = 3.368104 + 0.021748x_{2j}$$

Skim

$$\hat{y} = 3.04920 + 0.05796x_{3j}$$

We could have also had a more involved calculation of the above estimates using:

```
reduced <- lm(lnConc ~ source + percent + source:percent,
              data = pigs)
summary(reduced)

##
## Call:
## lm(formula = lnConc ~ source + percent + source:percent, data = pigs)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.195329 -0.043668 -0.001485  0.059798  0.169293
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   3.177909   0.085506  37.166 < 2e-16 ***
## source1      -0.061489   0.119140  -0.516   0.611
## source2       0.190195   0.121217   1.569   0.130
## percent       0.033603   0.006477   5.188 2.93e-05 ***
## source1:percent -0.012503   0.008723  -1.433   0.165
## source2:percent -0.011855   0.009277  -1.278   0.214
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1069 on 23 degrees of freedom
## Multiple R-squared:  0.7901, Adjusted R-squared:  0.7444
## F-statistic: 17.31 on 5 and 23 DF,  p-value: 3.968e-07
```

Where our “baseline” is for the “fish” source.

## Q2

Consider the plant density example discussed in slide set 6.

Source	Sum of Squares	DF
$x 1$	$\mathbf{y}^\top (\mathbf{P}_2 - \mathbf{P}_1) \mathbf{y}$	$2 - 1 = 1$
$x^2 1, x$	$\mathbf{y}^\top (\mathbf{P}_3 - \mathbf{P}_2) \mathbf{y}$	$3 - 2 = 1$
$x^3 1, x, x^2$	$\mathbf{y}^\top (\mathbf{P}_4 - \mathbf{P}_3) \mathbf{y}$	$4 - 3 = 1$
$x^4 1, x, x^2, x^3$	$\mathbf{y}^\top (\mathbf{P}_5 - \mathbf{P}_4) \mathbf{y}$	$5 - 4 = 1$
Error	$\mathbf{y}^\top (\mathbf{I} - \mathbf{P}_5) \mathbf{y}$	$15 - 5 = 10$
C. Total	$\mathbf{y}^\top (\mathbf{I} - \mathbf{P}_1) \mathbf{y}$	$15 - 1 = 14$

Figure 1: CocoMelon

a)

For each of the tests in the ANOVA table on slide 38, provide a vector  $c$  so that a test of

$$H_0 : c^T \beta = 0$$

would yield the same statistic and p-value as the ANOVA test. (You can use R to help you with the computations like we did on slides 45 and 46 of slide set 6.) Label these vectors  $c_1$ ,  $c_2$ ,  $c_3$ , and  $c_4$  for the linear, quadratic, cubic, and quartic tests, respectively.

```
proj <- function(x) {
  x %*% MASS::ginv(t(x) %*% x) %*% t(x)
}

# Had to dig through Canvas to find the supplemental R handout to find this. Smh
# Or, I guess not, there's a lot there to borrow for this problem
d <- read.delim("https://dnett.github.io/S510/PlantDensity.txt")
names(d) <- c("x", "y")

n <- nrow(d)
x <- (d$x - mean(d$x)) / 10
# Iteratively, we just need to add a higher order term to the prior design matrix
x1 <- matrix(1, nrow = n, ncol = 1)
x2 <- cbind(x1, x)
x3 <- cbind(x2, x^2)
x4 <- cbind(x3, x^3)
x5 <- matrix(model.matrix(~0 + factor(x)), nrow <- n)

p1 <- proj(x1)
p2 <- proj(x2)
p3 <- proj(x3)
p4 <- proj(x4)
```

```
p5 <- proj(x5)

# Contrasts
((p2-p1)%*%x5)[1,] * 5
```

```
## [1] 2 1 0 -1 -2
```

```
((p3-p2)%*%x5)[1,] * 7
```

```
## [1] 2 -1 -2 -1 2
```

```
((p4-p3)%*%x5)[1,] * 10
```

```
## [1] 1 -2 0 2 -1
```

```
((p5-p4)%*%x5)[1,] * 70
```

```
## [1] 1 -4 6 -4 1
```

From the above output, we have:

$$\mathbf{c}_1^\top = [2, 1, 0, -1, -2]$$

$$\mathbf{c}_2^\top = [2, -1, -2, -1, 2]$$

$$\mathbf{c}_3^\top = [1, -2, 0, 2, -1]$$

$$\mathbf{c}_4^\top = [1, -4, 6, -4, 1]$$

b)

Are  $c_1^T \beta$ ,  $c_2^T \beta$ ,  $c_3^T \beta$ , and  $c_4^T \beta$  contrasts? Explain.

All  $\mathbf{c}_i^T \beta$  are contrasts because:

$$\mathbf{c}_i^T \mathbf{1} = 0 \quad \text{for } i = 1, 2, 3, 4$$

Explicitly:

$$\mathbf{c}_1^T \mathbf{1} = 2 + 1 + 0 - 1 - 2 = 0$$

$$\mathbf{c}_2^T \mathbf{1} = 2 - 1 - 2 - 1 + 2 = 0$$

$$\mathbf{c}_3^T \mathbf{1} = 1 - 2 + 0 + 2 - 1 = 0$$

$$\mathbf{c}_4^T \mathbf{1} = 1 - 4 + 6 - 4 + 1 = 0$$

And not only are these contrasts, but they are orthogonal contrasts because they satisfy:

$$\mathbf{c}_i^T \mathbf{c}_j = 0, \quad \text{for } i \neq j$$

c)

Are  $c_1^T \beta$ ,  $c_2^T \beta$ ,  $c_3^T \beta$ , and  $c_4^T \beta$  orthogonal? Explain.

Preempted this question a bit at the end of part b), but yes, these contrasts are orthogonal, and here is more detail on why beyond the expression they satisfy.

So, the initial expression to check is:

$$\mathbf{c}_i^\top \mathbf{c}_j = 0, \quad \text{for } i \neq j$$

Explicitly:

4 choose 2 cases to check, 6 total cases:

$$\mathbf{c}_1^\top \mathbf{c}_2 = (2)(2) + (1)(-1) + (0)(-2) + (-1)(-1) + (-2)(2) = 4 - 1 + 0 + 1 - 4 = 0$$

$$\mathbf{c}_1^\top \mathbf{c}_3 = (2)(1) + (1)(-2) + (0)(0) + (-1)(2) + (-2)(-1) = 2 - 2 + 0 - 2 + 2 = 0$$

$$\mathbf{c}_1^\top \mathbf{c}_4 = (2)(1) + (1)(-4) + (0)(6) + (-1)(-4) + (-2)(1) = 2 - 4 + 0 + 4 - 2 = 0$$

$$\mathbf{c}_2^\top \mathbf{c}_3 = (2)(1) + (-1)(-2) + (-2)(0) + (-1)(2) + (2)(-1) = 2 + 2 + 0 - 2 - 2 = 0$$

$$\mathbf{c}_2^\top \mathbf{c}_4 = (2)(1) + (-1)(-4) + (-2)(6) + (-1)(-4) + (2)(1) = 2 + 4 - 12 + 4 + 2 = 0$$

$$\mathbf{c}_3^\top \mathbf{c}_4 = (1)(1) + (-2)(-4) + (0)(6) + (2)(-4) + (-1)(1) = 1 + 8 + 0 - 8 - 1 = 0$$

Note: The above is a simplification that works for the example given because of the design matrix used in this problem.

The “base” definition we reference for testing whether something is orthogonal is: Any two estimable linear combinations  $c_i^T \beta$  and  $c_j^T \beta$  are orthogonal if and only if:

$$\mathbf{c}_i^T (\mathbf{X}^T \mathbf{X})^{-} \mathbf{c}_j = 0 \quad \text{for } i \neq j$$

We use the full design matrix (`x5` from above) for the calculations. Since  $\mathbf{X}^T \mathbf{X}$  is full rank, we can compute its inverse directly rather than using a generalized inverse. Specifically, because  $\mathbf{X}^T \mathbf{X}$  is a scalar multiple of the identity matrix, we ensure it is invertible, meaning its inverse is unique and can be computed using the `solve` function in R.

Evaluating the above, we get the following matrices:

```
require(MASS)
```

```
## Loading required package: MASS
```

```
##
```

```
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':
##
##      select
```

```
x5
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    0    0    0    0
## [2,]    1    0    0    0    0
## [3,]    1    0    0    0    0
## [4,]    0    1    0    0    0
## [5,]    0    1    0    0    0
## [6,]    0    1    0    0    0
## [7,]    0    0    1    0    0
## [8,]    0    0    1    0    0
## [9,]    0    0    1    0    0
## [10,]   0    0    0    1    0
## [11,]   0    0    0    1    0
## [12,]   0    0    0    1    0
## [13,]   0    0    0    0    1
## [14,]   0    0    0    0    1
## [15,]   0    0    0    0    1
```

```
t(x5) %*% x5
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    3    0    0    0    0
## [2,]    0    3    0    0    0
## [3,]    0    0    3    0    0
## [4,]    0    0    0    3    0
## [5,]    0    0    0    0    3
```

```
fractions(solve(t(x5) %*% x5))
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,] 1/3    0    0    0    0
## [2,] 0 1/3    0    0    0
## [3,] 0    0 1/3    0    0
## [4,] 0    0    0 1/3    0
## [5,] 0    0    0    0 1/3
```

Thus, in this case,

$$\mathbf{c}_i^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{c}_j = \mathbf{c}_i^T \mathbf{c}_j / 3$$

so that linear combinations  $\mathbf{c}_i^T \boldsymbol{\beta}$  and  $\mathbf{c}_j^T \boldsymbol{\beta}$  are orthogonal if and only if  $\mathbf{c}_i^T \mathbf{c}_j = 0$ . (Multiplying each side by 3 for simplicity).

The rest of the results (that our contrasts are orthogonal) follow from the explicit derivations.

### Q3

Suppose  $\mathbf{H}$  is a symmetric matrix. Prove that  $H$  is nonnegative definite if and only if all its eigenvalues are nonnegative. (If you wish, you may use the Spectral Decomposition Theorem in your proof.)

I do want to use Spectral Decomposition, thank you!

#### Spectral Decomposition Theorem:

For  $\mathbf{H}$  is a symmetric matrix, then:

$$\mathbf{H} = \mathbf{P}\mathbf{\Lambda}\mathbf{P}^\top = \sum_{i=1}^n \lambda_i \mathbf{p}_i \mathbf{p}_i^\top$$

where  $\mathbf{p}_i$  are orthonormal eigenvectors of  $\mathbf{H}$ .

The general approach for solving an iff proof is to prove both directions hold. To that end:

#### Direction 1

Assume we have  $\mathbf{H}$ , a symmetric, nonnegative definite matrix.

By definition,  $\mathbf{H}$  being nonnegative definite, means the following holds:

$$\mathbf{p}_i^\top \mathbf{H} \mathbf{p}_i \geq 0$$

for any  $\mathbf{p}_i$ , where  $i = 1, \dots, n$ .

Unfurling  $\mathbf{H}$ , we may rewrite:

$$\mathbf{p}_i^\top \mathbf{H} \mathbf{p}_i = \mathbf{p}_i^\top \left( \sum_{j=1}^n \lambda_j \mathbf{p}_j \mathbf{p}_j^\top \right) \mathbf{p}_i = \sum_{j=1}^n \lambda_j \mathbf{p}_i^\top \mathbf{p}_j \mathbf{p}_j^\top \mathbf{p}_i$$

By Special Decomposition, we know  $\mathbf{p}_i$  are orthonormal, meaning:

$$\mathbf{p}_i^\top \mathbf{p}_j = 0 \quad \text{for all } i \neq j, \quad \text{and} \quad \mathbf{p}_i^\top \mathbf{p}_i = 1$$

Allowing us to further simplify:

$$\mathbf{p}_i^\top \mathbf{H} \mathbf{p}_i = \lambda_i \mathbf{p}_i^\top \mathbf{p}_i \mathbf{p}_i^\top \mathbf{p}_i = \lambda_i$$

And because  $\mathbf{H}$  is nonnegative definite, it then follows that:

$$\lambda_i \geq 0, \quad \forall i = 1, \dots, n$$

Direction one complete!



## Direction 2

We start by assuming that for a symmetric matrix  $\mathbf{H}$ , all its eigenvalues are nonnegative.

Given our assumption, we may restate as  $\lambda_i \geq 0$  for  $i = 1, \dots, n$ .

By the Spectral Decomposition Theorem:

$$\mathbf{H} = \mathbf{P} \text{diag}(\lambda_1, \dots, \lambda_n) \mathbf{P}^\top$$

where  $\mathbf{P} = [\mathbf{p}_1, \dots, \mathbf{p}_n]$  and:

$$\mathbf{P} \mathbf{P}^\top = \mathbf{P}^\top \mathbf{P} = \mathbf{I} \rightarrow \mathbf{p}_j^\top \mathbf{p}_j = 1$$

For  $j = 1, \dots, n$ , Expressing  $\mathbf{y}$  in terms of the eigenvectors of  $\mathbf{H}$ , we have:

$$\mathbf{y} = \sum_{j=1}^n x_j \mathbf{p}_j \rightarrow x_j = \mathbf{p}_j^\top \mathbf{y}$$

By Spectral Decomposition, and then simplifying:

$$\mathbf{y}^\top \mathbf{H} \mathbf{y} = \mathbf{y}^\top \mathbf{P} \Lambda \mathbf{P}^\top \mathbf{y} = \mathbf{y}^\top \left( \sum_{j=1}^n \lambda_j \mathbf{p}_j \mathbf{p}_j^\top \right) \mathbf{y} = \sum_{j=1}^n \lambda_j \mathbf{y}^\top \mathbf{p}_j \mathbf{p}_j^\top \mathbf{y} = \sum_{j=1}^n \lambda_j x_j^2$$

Individually, we know  $x_j^2$  is non-negative. And as given, we know  $\lambda_j$  is non-negative. Taken together, we know that each term  $\lambda_j x_j^2$  is non-negative, such that their sum is non-negative as well. This means we have shown:

$$\mathbf{y}^\top \mathbf{H} \mathbf{y} \geq 0$$

Making the matrix  $\mathbf{H}$  nonnegative definite.

Second direction complete!

## If

Taken together, for  $H$  is a symmetric matrix,  $H$  is nonnegative definite  $\iff$  all its eigenvalues are nonnegative.