

Assignment 7

Sam Olson

The Assignment

Your assignment is to conduct a Bayesian analysis of relative risk for these data. You will do so using both the Multinomial probability model and the Poisson probability model.

1.

Using prior distributions and prior parameter values as given in Section 4 (Bayesian Analysis), derive the joint posterior distributions for $\boldsymbol{\theta} = (\theta_{11}, \theta_{10}, \theta_{01}, \theta_{00})$ in the multinomial model, and $\boldsymbol{\lambda} = (\lambda_{11}, \lambda_{10}, \lambda_{01}, \lambda_{00})$ in the Poisson model.

Answer

Multinomial–Dirichlet Multinomial likelihood (with $N = \sum_{ij} y_{ij}$, $\sum_{ij} \theta_{ij} = 1$):

$$L(\boldsymbol{\theta} \mid \mathbf{y}) = \Pr(\mathbf{Y} = \mathbf{y} \mid \boldsymbol{\theta}) = \frac{N!}{\prod_{ij} y_{ij}!} \prod_{ij} \theta_{ij}^{y_{ij}}$$

For the priors:

Dirichlet prior on cell probabilities $\boldsymbol{\theta}$ (with support on a simplex, i.e., $\sum_{i,j} \theta_{ij} = 1$):

$$\boldsymbol{\theta} \sim \text{Dirichlet}(\gamma_{11}, \gamma_{10}, \gamma_{01}, \gamma_{00}), \quad p(\boldsymbol{\theta}) = \frac{1}{B(\boldsymbol{\gamma})} \prod_{ij} \theta_{ij}^{\gamma_{ij}-1}$$

And

$$B(\boldsymbol{\gamma}) = \frac{\prod_{ij} \Gamma(\gamma_{ij})}{\Gamma(\sum_{ij} \gamma_{ij})},$$

As given, $(\gamma_{11}, \gamma_{10}, \gamma_{01}, \gamma_{00}) = (0.10, 0.75, 0.10, 0.75)$.

Taken together, we have the (proportional) posterior derivations

Multinomial–Dirichlet:

$$\begin{aligned} p(\boldsymbol{\theta} \mid \mathbf{y}) &\propto L(\boldsymbol{\theta} \mid \mathbf{y}), p(\boldsymbol{\theta}) \\ &\propto \left(\prod_{ij} \theta_{ij}^{y_{ij}} \right) \left(\prod_{ij} \theta_{ij}^{\gamma_{ij}-1} \right) \\ &\propto \prod_{ij} \theta_{ij}^{y_{ij} + \gamma_{ij} - 1} \end{aligned}$$

For $\sum_{ij} \theta_{ij} = 1$

After collecting exponents and normalizing over the simplex (by the Dirichlet Beta function), conjugacy gives

$$\boldsymbol{\theta} \mid \mathbf{y} \sim \text{Dirichlet}(y_{11} + \gamma_{11}, y_{10} + \gamma_{10}, y_{01} + \gamma_{01}, y_{00} + \gamma_{00})$$

Or, explicitly with the given values, the posterior is:

$$\boldsymbol{\theta} \mid \mathbf{y} \sim \text{Dirichlet}(15.10, 50.75, 3.10, 84.75)$$

Poisson–Gamma (Independent) Poisson likelihoods:

$$Y_{ij} \mid \lambda_{ij} \stackrel{\text{ind}}{\sim} \text{Poisson}(\lambda_{ij}) \Rightarrow L(\boldsymbol{\lambda} \mid \mathbf{y}) = \prod_{ij} \frac{e^{-\lambda_{ij}} \lambda_{ij}^{y_{ij}}}{y_{ij}!}$$

For the priors:

Gamma priors (assuming shape–rate parametrization) on Poisson means:

$$\lambda_{ij} \sim \text{Gamma}(\alpha_{ij}, \beta_{ij}), \quad p(\lambda_{ij}) = \frac{\beta_{ij}^{\alpha_{ij}}}{\Gamma(\alpha_{ij})} \lambda_{ij}^{\alpha_{ij}-1} e^{-\beta_{ij} \lambda_{ij}}$$

Given $(\alpha, \beta)_{11,10,01,00} = (0.5, 5), (3, 4), (0.5, 5), (3, 4)$.

Taken together, we have the (proportional) posterior derivations

$$\begin{aligned} p(\lambda_{ij} \mid y_{ij}) &\propto e^{-\lambda_{ij}} \lambda_{ij}^{y_{ij}} \times \lambda_{ij}^{\alpha_{ij}-1} e^{-\beta_{ij} \lambda_{ij}} \\ &\propto \lambda_{ij}^{\alpha_{ij}+y_{ij}-1} \exp[-(\beta_{ij} + 1)\lambda_{ij}] \end{aligned}$$

Recognizing this kernel as a Gamma density, conjugacy gives us posterior(s) of the form:

$$\lambda_{ij} \mid y_{ij} \sim \text{Gamma}(\alpha_{ij} + y_{ij}, \beta_{ij} + 1)$$

Or, explicitly with the given values, the posterior distributions are:

$$\lambda_{11} \mid y_{11} \sim \text{Gamma}(15.5, 6)$$

$$\lambda_{10} \mid y_{10} \sim \text{Gamma}(53, 5)$$

$$\lambda_{01} \mid y_{01} \sim \text{Gamma}(3.5, 6)$$

$$\lambda_{00} \mid y_{00} \sim \text{Gamma}(87, 5)$$

2.

Using the data of Table 2 give posterior expected values for θ and λ , and the cell probabilities that would correspond to these means.

Answer

We have observed counts $y_{11} = 15$, $y_{10} = 50$, $y_{01} = 3$, $y_{00} = 84$

So our total N is given by: $N = 152$

Let $G = \sum \gamma_{ij} = 1.70$,

Then: $N + G = 153.70$

Multinomial–Dirichlet Using properties of known distributions, for Dirichlet:

$$E[\theta_{ij}] = \gamma_{ij}/G$$

And

$$\text{Var}(\theta_{ij}) = \frac{\gamma_{ij}(G - \gamma_{ij})}{G^2(G + 1)}.$$

So:

$$E[\theta_{ij} | y] = \frac{y_{ij} + \gamma_{ij}}{N + G}$$

Evaluating,

$$\begin{aligned} E[\theta_{11} | y] &= \frac{15 + 0.10}{153.70} = 0.0982, \\ E[\theta_{10} | y] &= \frac{50 + 0.75}{153.70} = 0.3302, \\ E[\theta_{01} | y] &= \frac{3 + 0.10}{153.70} = 0.0202, \\ E[\theta_{00} | y] &= \frac{84 + 0.75}{153.70} = 0.5514 \end{aligned}$$

Note: We don't need to normalize the above, as we have $\sum_{ij} \theta_{ij} = 0.098 + 0.330 + 0.020 + 0.551 = 1$ (well, ok: 0.999 but that's just due to rounding!)

Poisson–Gamma Again, for a known Gamma distribution (as parametrized):

$$E[\lambda_{ij}] = \alpha_{ij}/\beta_{ij},$$

And

$$\text{Var}(\lambda_{ij}) = \alpha_{ij}/\beta_{ij}^2.$$

So, we have:

$$E[\lambda_{ij} | y] = \frac{\alpha_{ij} + y_{ij}}{\beta_{ij} + 1}$$

So

$$\begin{aligned}
E[\lambda_{11}] &= \frac{15.5}{6} = 2.5833 \\
E[\lambda_{10}] &= \frac{53}{5} = 10.6 \\
E[\lambda_{01}] &= \frac{3.5}{6} = 0.5833 \\
E[\lambda_{00}] &= \frac{87}{5} = 17.4
\end{aligned}$$

To compare on the probability scale, we need to normalize:

$$\tilde{\theta}_{ij} = \frac{E[\lambda_{ij} \mid y]}{\sum_{kl} E[\lambda_{kl} \mid y]} = (0.0829, 0.3401, 0.0187, 0.5583)$$

Note: I believe the phrasing of this question is asking for the unnormalized posterior expectations ($E[\lambda_{ij}] = (\alpha_{ij} + y_{ij})/(\beta_{ij} + 1)$), which were first provided in this section's write up. However, I also provide the normalized quantities $\tilde{\theta}_{ij}$ because I am using those quantities for plug-in RR calculations in Question 3 (and also because it's somewhat unclear which/if both are being explicitly asked for).

3.

Evaluate relative risk RR at the expected values of the posterior distributions you derived in exercise 1 and compare to the observed or sample-based RR based only on the observed cell counts. Briefly explain why this is not the posterior expected value of RR (although we hope it is not far off).

Answer

The relative risk is defined as in expression (1) of the assignment as:

$$RR = \frac{\Pr(D | R)}{\Pr(D | R^c)} = \frac{\theta_{11}(\theta_{01} + \theta_{00})}{\theta_{01}(\theta_{11} + \theta_{10})}$$

With the sample version of RR as defined in expression (2) of the assignment, we have:

$$\widehat{RR} = \frac{y_{11}(y_{01} + y_{00})}{y_{01}(y_{11} + y_{10})} = \frac{15(3 + 84)}{3(15 + 50)} \approx 6.69$$

Using the posterior “plug-in” values, we have:

Multinomial–Dirichlet:

$$RR(E[\theta]) = \frac{0.0982(0.0202 + 0.5514)}{0.0202(0.0982 + 0.3302)} \approx 6.50$$

And

Poisson–Gamma:

$$RR(\tilde{\theta}) = \frac{0.0829(0.0187 + 0.5583)}{0.0187(0.0829 + 0.3401)} \approx 6.04$$

Thoughts All three values are close, but not exactly the same. Although we are using the posterior distributions of the components of RR (the θ_{ij} ’s or λ_{ij} ’s), the “plug-in” value RR is not the same as the posterior mean $E[RR | y]$, which would require integrating the function $RR = g(\theta)$ over its full posterior distribution. This difference primarily stems from the nonlinearity of the function RR , which means Jensen’s inequality applies: Expectations do not generally “pass through” nonlinear transformations such that they ensure equality.

To be specific, the function g for RR is given by

$$g(\theta) = \frac{\theta_{11}(\theta_{01} + \theta_{00})}{\theta_{01}(\theta_{11} + \theta_{10})}$$

and because g is nonlinear in θ ,

$$E[g(\theta) | y] \neq g(E[\theta | y])$$

Though this might be repetitive of the above, taking an expectation and then applying a nonlinear function (like this ratio for RR) does not yield the same result as applying the function first and then averaging. The difference reflects the variability in the posterior distribution of $\theta | y$ (and also the specific nonlinear form of the function g). Finally, the observed sample estimate \widehat{RR} also differs slightly because it is based only on the counts, while the Bayesian “plug-in” values incorporate both the data likelihood and prior information.

Note: For the Poisson–Gamma model, one could equivalently parameterize using the counts (λ_{ij}) rather than the normalized probabilities ($\theta_{ij} = \lambda_{ij}/\lambda_N$), as noted in the assignment. The choice of probabilities over counts was done here for convenience, and not due to any willful decision. However, either parametrization are functionally identical for the purpose of this problem, as relative risk depends only on ratios that cancel the overall scale of (λ). That, and: The goal here is to compare results *between* the Poisson–Gamma and Multinomial–Dirichlet (and the sample-based estimates), and not between two equivalent parameterizations of the Poisson–Gamma model. And, as noted in the assignment, these two representations yield “*computationally the same results*”.

4.

Find the posterior distributions of RR under both the multinomial and Poisson models. Give summary and 95% credible intervals. Produce some type of a graphical display of the posterior distributions. How do the results compare for the multinomial and Poisson models? How do the results compare to the initial sample version of RR in expression (2)? What would you conclude about the relation between time since vaccination and the chances of getting chickenpox in this population of school children? For example, what is your posterior probability that the risk factor of having more than 5 years since vaccination is positively related to the chance of contracting chickenpox in these students?

Answer

For the Multinomial-Dirichlet, we have:

$$\boldsymbol{\theta} \mid \mathbf{y} \sim \text{Dirichlet}(y_{11} + \gamma_{11}, y_{10} + \gamma_{10}, y_{01} + \gamma_{01}, y_{00} + \gamma_{00})$$

And for the Poisson-Gamma, we have:

$$\lambda_{ij} \mid \mathbf{y} \sim \text{Gamma}(\alpha_{ij} + y_{ij}, \beta_{ij} + 1)$$

Which requires normalization via:

$$p_{ij} = \frac{\lambda_{ij}}{\sum_{kl} \lambda_{kl}}$$

The form of RR is then given by:

For the Multinomial-Dirichlet:

$$RR(\boldsymbol{\theta}) = \frac{\theta_{11}(\theta_{01} + \theta_{00})}{\theta_{01}(\theta_{11} + \theta_{10})}$$

And for the Poisson-Gamma:

$$RR(\mathbf{p}) = \frac{p_{11}(p_{01} + p_{00})}{p_{01}(p_{11} + p_{10})}$$

The procedure then, using the above, is to generate posterior draws for the parameters using `rdirichlet` (for the Multinomial-Dirichlet) and `rgamma` (for the Poisson-Gamma model). For each draw, an RR is computed using the corresponding sampled probabilities (and for each model). The resulting collection of RR values then is used as our “empirical” posterior distribution of RR. It is the distribution of these RR values then that are summarized; particularly, the credible interval is determined by the empirical 2.5% and 97.5% quantiles of the (simulated) RR distribution.

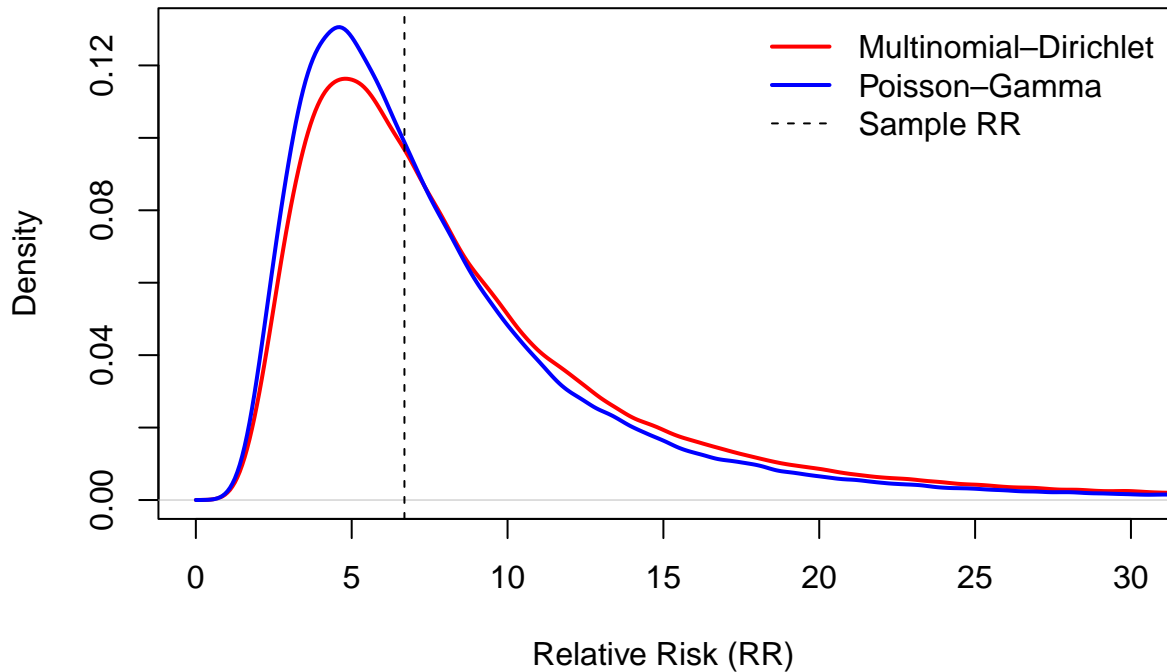
Special thanks to Kaiser for noting that we don’t need to do any special (rejection) sampling for this problem.

Note: Similar to the note in the preceding problem, we could also generate posterior samples directly from the Poisson–Gamma counts (λ_{ij}) instead of from the normalized probabilities ($p_{ij} = \lambda_{ij} / \sum_{kl} \lambda_{kl}$). However, and again, both approaches are functionally equivalent for computing relative risk, since the overall scale of (λ) cancels in the RR ratio. And, again, the purpose here is to compare the posterior behavior of RR across the Multinomial–Dirichlet and Poisson–Gamma formulations, and not to contrast two equivalent ways of sampling *within* the Poisson–Gamma.

Table 1: Posterior summaries for RR under two Bayesian models

Model	Posterior Summary						
	Mean	Median	SD	2.5%	97.5%	$\Pr(RR > 1)$	$\Pr(RR > 5)$
Multinomial–Dirichlet	9.50	7.10	9.26	2.42	31.18	1	0.72
Poisson–Gamma	8.38	6.51	7.08	2.28	25.76	1	0.68

Posterior Distributions of RR



Thoughts Both posterior distributions of RR place nearly all their (respective) mass(es) above 1 and are “concentrated” (are most dense) in the range of values between 4 to 7. Relative to the sample estimate $\widehat{RR} \approx 6.69$, the “plug-in” estimates show mild shrinkage ($RR(E[\theta]) \approx 6.50$ and $RR(\tilde{\theta}) \approx 6.04$).

Most importantly for our interpretation, the posterior probability that $RR > 1$ is effectively 1 under both models:

$$\Pr(RR > 1 \mid \text{Multinomial–Dirichlet}) \approx 0.99996, \quad \Pr(RR > 1 \mid \text{Poisson–Gamma}) \approx 0.99990.$$

Thus, there is very strong posterior evidence that “> 5 years since vaccination” is associated with a higher probability of contracting/“getting” chickenpox in the student population. The 95% credible intervals are somewhat wide on the upper tail (reflecting uncertainty for large RR values) but largely overlap across the two models and include the sample \widehat{RR} , indicating consistent conclusions across the two modeling frameworks (Multinomial-Dirichlet and Poisson-Gamma).