

# HW3

Sam Olson

## Outline

- Q1: g2g
- Q2: g2g
- Q3:
- Q4:
- Q5: little weird at parts, with  $\propto$

### 1.

Suppose  $X_1, \dots, X_n$  are iid Bernoulli( $p$ ),  $0 < p < 1$ .

a)

Find the information number  $I_n(p)$  and make a rough sketch of  $I_n(p)$  as a function of  $p \in (0, 1)$ .

Given that  $X_1, \dots, X_n$  are i.i.d. Bernoulli( $p$ ), the likelihood function is:

$$L(p) = \prod_{i=1}^n p^{X_i} (1-p)^{1-X_i}$$

Taking the log-likelihood,

$$\log(L(p)) = \sum_{i=1}^n [X_i \log p + (1 - X_i) \log(1 - p)]$$

The first derivative is:

$$\log(L(p))' = \sum_{i=1}^n \left[ \frac{X_i}{p} - \frac{1 - X_i}{1 - p} \right] = \sum_{i=1}^n \frac{X_i - p}{p(1 - p)}$$

The Fisher information is:

$$I_n(p) = -E[\log(L(p))'']$$

Computing the second derivative:

$$\log(L(p))'' = \sum_{i=1}^n \left[ -\frac{X_i}{p^2} - \frac{1 - X_i}{(1 - p)^2} \right]$$

Taking expectation:

$$E[\log(L(p))''] = \sum_{i=1}^n \left[ -\frac{E[X_i]}{p^2} - \frac{E[1 - X_i]}{(1-p)^2} \right]$$

Given we know the distribution of the random variables, we know  $E[X_i] = p$  and  $E[1 - X_i] = 1 - p$ . This allows us to simplify the expression:

$$E[\log(L(p))''] = \sum_{i=1}^n \left[ -\frac{p}{p^2} - \frac{1-p}{(1-p)^2} \right] = \sum_{i=1}^n \left[ -\frac{1}{p} - \frac{1}{1-p} \right] = -n \left[ \frac{1}{p} + \frac{1}{1-p} \right]$$

Noting linearity of Fisher information:

$$I_n(p) = n \left[ \frac{1}{p} + \frac{1}{1-p} \right]$$


**Sketch**

```
# functional form
fisher_info <- function(p, n) {
  return(n * (1/p + 1/(1 - p)))
}

# setup
p_values <- seq(0.01, 0.99, length.out = 100)
n <- 1
I_values <- fisher_info(p_values, n)

# plot
plot(x = p_values,
     y = I_values,
     type = "l",
     col = "black", lwd = 2,
     xlab = "p", ylab = "Fisher Information",
     main = "Fisher Information for Bernoulli(p)")
abline(v = 0.5, lty = 2, col = "gray")
```

## Fisher Information for Bernoulli(p)



b)

Find the value of  $p \in (0, 1)$  for which  $I_n(p)$  is minimal. (This value of  $p$  corresponds to the “hardest” case for estimating  $p$ . That is, when data are generated under this value of  $p$  from the model, the variance of an UE of  $p$  is potentially largest.)

To find the value of  $p$  that minimizes the Fisher information  $I_n(p)$ , we use the functional form of the Fisher Information:

$$I_n(p) = n \left[ \frac{1}{p} + \frac{1}{1-p} \right]$$

Differentiating  $I_n(p)$  with respect to  $p$ , and setting equal to zero:

$$I_n(p)' = n \left[ -\frac{1}{p^2} + \frac{1}{(1-p)^2} \right] = -\frac{1}{p^2} + \frac{1}{(1-p)^2} = 0$$

This gives us the expression:

$$\frac{1}{p^2} = \frac{1}{(1-p)^2}$$

Taking square roots:

$$\frac{1}{p} = \frac{1}{1-p} \rightarrow p = 1-p \rightarrow p = \frac{1}{2}$$

To ensure this is a maximum, we also check whether the second derivative is positive (since we are minimizing and not maximizing) at  $\frac{1}{2}$ :

$$I_n(p)' = n \left[ \frac{2}{p^3} + \frac{2}{(1-p)^3} \right]$$

$$I_n \left( \frac{1}{2} \right)'' = n \left[ \frac{2}{(1/2)^3} + \frac{2}{(1/2)^3} \right] = n \left[ \frac{2}{1/8} + \frac{2}{1/8} \right] = n [16 + 16] = 32n > 0$$

So this is in fact a minimum, hence the Fisher information is minimized at:

$$p = \frac{1}{2}$$

c)

Show that  $\hat{X}_n = \sum_{i=1}^n X_i/n$  is the UMVUE of  $p$ .

Note to self: Uniformly Minimum Variance Unbiased Estimator (UMVUE)

We start by checking if  $\hat{X}_n$  is an unbiased estimator of  $p$ :

$$E[\hat{X}_n] = E \left[ \frac{1}{n} \sum_{i=1}^n X_i \right] = \frac{1}{n} \sum_{i=1}^n E[X_i] E[\hat{X}_n] = \frac{1}{n} \cdot np = p$$

$$Bias(\bar{X}_n) = E[\hat{X}_n] - E[X] = p - p = 0$$

So  $\hat{X}_n$  is an unbiased estimator of  $p$ .

Now as far as the “Uniformly Minimum Variance” part of the question:

Note again the Fisher Information formula we’ve found:

$$I_n(p) = \frac{np}{p^2} + \frac{n(1-p)}{(1-p)^2} = \frac{n}{p(1-p)}$$

By the definition, the Cramér-Rao Lower Bound, for any unbiased estimator  $T$  of  $p$ :

$$\text{Var}_p(T) \geq \frac{(\gamma'(p))^2}{I_n(p)}$$

Here, we are estimating  $\gamma(p) = p$ , so  $\gamma'(p) = 1$ . Therefore:

$$\text{Var}_p(T) \geq \frac{1^2}{I_n(p)} = \frac{p(1-p)}{n}$$

We compute the variance of  $\hat{X}_n = S_n/n$ :

$$E[\hat{X}_n] = E \left[ \frac{S_n}{n} \right] = \frac{1}{n} E[S_n] = \frac{np}{n} = p$$

$$\text{Var}(\hat{X}_n) = \text{Var} \left( \frac{S_n}{n} \right) = \frac{1}{n^2} \text{Var}(S_n)$$

Since  $S_n \sim \text{Binomial}(n, p)$ , we know:

$$\text{Var}(S_n) = np(1 - p)$$

Thus:

$$\text{Var}(\hat{X}_n) = \frac{np(1 - p)}{n^2} = \frac{p(1 - p)}{n}$$

Comparing with the CRLB:

$$\text{Var}(\hat{X}_n) = \frac{p(1 - p)}{n} = \frac{1}{I_n(p)}$$

Since  $\hat{X}_n$  attains the bound, it is an efficient estimator.

Since  $\hat{X}_n$  is unbiased and attains the CRLB, it is the UMVUE.

## 2.

Suppose that the random variables  $Y_1, \dots, Y_n$  satisfy

$$Y_i = \beta x_i + \varepsilon_i, \quad i = 1, \dots, n$$

where  $x_1, \dots, x_n$  are fixed constants and  $\varepsilon_1, \dots, \varepsilon_n$  are iid  $N(0, \sigma^2)$ ; here we assume  $\sigma^2 > 0$  is known.

a)

Find the MLE of  $\beta$ .

To find the Maximum Likelihood Estimator (MLE) of  $\beta$ , we first write the likelihood function.

Since  $Y_i = \beta x_i + \varepsilon_i$  with  $\varepsilon_i \sim N(0, \sigma^2)$ , we have:

$$Y_i \sim N(\beta x_i, \sigma^2)$$

Thus, the joint density function of  $Y_1, \dots, Y_n$  is:

$$L(\beta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(Y_i - \beta x_i)^2}{2\sigma^2}\right)$$

Taking the log-likelihood:

$$\log(L(\beta)) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (Y_i - \beta x_i)^2$$

To find the MLE of  $\beta$ , we take the derivative with respect to  $\beta$  and set to zero:

$$\frac{d}{d\beta} \log(L(\beta)) = -\frac{1}{2\sigma^2} \cdot (-2) \sum_{i=1}^n x_i (Y_i - \beta x_i) = \frac{1}{\sigma^2} \sum_{i=1}^n x_i (Y_i - \beta x_i) \rightarrow \sum_{i=1}^n x_i Y_i - \beta \sum_{i=1}^n x_i^2 = 0$$

Solving for  $\beta$ , we get our MLE of  $\beta$  as::

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i Y_i}{\sum_{i=1}^n x_i^2}$$

To ensure this is a maximum, we take the second derivative and see if it is negative:

$$\log(L(\beta))'' = -\sum_{i=1}^n x_i^2 < 0$$

So this is in fact the maximum.

b)

Find the distribution of the MLE.

From part a), the MLE of  $\beta$  is:

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i Y_i}{\sum_{i=1}^n x_i^2}$$

To determine the distribution of  $\hat{\beta}$ , determine its expectation and variance, noting that since  $\hat{\beta}$  is a linear combination of the normal random variables  $\varepsilon_i$ , it follows that  $\hat{\beta}$  itself is normally distributed.

That being said, given  $Y_i = \beta x_i + \varepsilon_i$ , we may write:

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i(\beta x_i + \varepsilon_i)}{\sum_{i=1}^n x_i^2} = \frac{\beta \sum_{i=1}^n x_i^2 + \sum_{i=1}^n x_i \varepsilon_i}{\sum_{i=1}^n x_i^2}$$

Taking the expectation, noting our data is treated as “fixed”, we may write:

$$E[\hat{\beta}] = \frac{\beta \sum_{i=1}^n x_i^2 + \sum_{i=1}^n x_i E[\varepsilon_i]}{\sum_{i=1}^n x_i^2} = \frac{\beta \sum_{i=1}^n x_i^2}{\sum_{i=1}^n x_i^2} = \beta$$

Noting  $E[\varepsilon_i] = 0$

Because  $E[\hat{\beta}] = \beta$ , it has zero bias and  $\hat{\beta}$  is an unbiased estimator of  $\beta$ . Not needed for the distribution, but will need this note for later.

Let us then analyze the variance. We start again with definitions:

$$\text{Var}(\hat{\beta}) = \text{Var}\left(\beta + \frac{\sum_{i=1}^n x_i \varepsilon_i}{\sum_{i=1}^n x_i^2}\right) = \text{Var}(\beta) + \text{Var}\left(\frac{\sum_{i=1}^n x_i \varepsilon_i}{\sum_{i=1}^n x_i^2}\right)$$

Simplifying:

$$\text{Var}(\hat{\beta}) = \text{Var}\left(\frac{\sum_{i=1}^n x_i \varepsilon_i}{\sum_{i=1}^n x_i^2}\right) = \text{Var}\left(\sum_{i=1}^n \frac{(x_i^2 \sigma^2)}{(x_i^2)^2}\right) = \frac{\sigma^2}{\sum_{i=1}^n x_i^2}$$

We thus conclude:

$$\hat{\beta} \sim N\left(\beta, \frac{\sigma^2}{\sum_{i=1}^n x_i^2}\right)$$

c)

Find the CRLB for estimating  $\beta$ . (Hint: you'll have to work with the joint distribution  $f(y_1, \dots, y_n | \beta)$  directly, since  $Y_1, \dots, Y_n$  are not iid.)

To find the CRLB, we first calculate the Fisher information.

Note the joint density:

$$f(Y_1, \dots, Y_n | \beta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(Y_i - \beta x_i)^2}{2\sigma^2}\right)$$

Taking the log-likelihood:

$$\log(L(\beta)) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (Y_i - \beta x_i)^2$$

We take the derivative:

$$\log(L(\beta))' = -\frac{1}{2\sigma^2} \cdot (-2) \sum_{i=1}^n x_i (Y_i - \beta x_i) = \frac{1}{\sigma^2} \sum_{i=1}^n x_i (Y_i - \beta x_i)$$

The Fisher information is then:

$$I(\beta) = -E[\log(L(\beta))''] = -E \left[ -\frac{1}{\sigma^2} \sum_{i=1}^n x_i^2 \right] = \frac{1}{\sigma^2} \sum_{i=1}^n x_i^2$$

We then have what we need to calculate the CRLB using the information we've gathered.

The CRLB is:

$$\frac{1}{I(\beta)} = \frac{1}{\frac{1}{\sigma^2} \sum_{i=1}^n x_i^2} = \frac{\sigma^2}{\sum_{i=1}^n x_i^2}$$

d)

Show the MLE is the UMVUE of  $\beta$ .

Now we just need to compare the variance of our MLE of  $\beta$  to the value calculated in part c). To that end:

We have already calculated the expectation of  $\hat{\beta}_{MLE}$ , which is  $\beta$ , so via Bias calculation:

$$\text{Bias}(\hat{\beta}_{MLE}) = E[\hat{\beta}_{MLE}] - \beta = \beta - \beta = 0$$

Hence it is unbiased. We then just need to determine if our MLE attains the CRLB. If so, then the MLE is the UMVUE.

Recall the variance of the MLE:

$$\frac{\sigma^2}{\sum_{i=1}^n x_i^2}$$

And the CRLB:

$$\frac{\sigma^2}{\sum_{i=1}^n x_i^2}$$

These are one and the same! So we do indeed satisfy:

$$\text{Var}(\hat{\beta}_{MLE}) = CRLB$$

Such that the MLE is the UMVUE.



### 3.

Suppose  $X_1, \dots, X_n$  are iid normal  $N(0, 1)$ , where  $\theta \in \mathbb{R}$ . It turns out that  $T = (\bar{X}_n)^2 - n^{-1}$  is the UMVUE of  $\gamma(\theta) = \theta^2$ . (We can show this later in the course; our goal here is to show that the UMVUE can exist without obtaining the CRLB.)

a)

Show  $T$  is an UE of  $\gamma(\theta) = \theta^2$  and find the variance  $\text{Var}_\theta(T)$  of  $T$ . (Note  $Z = \sqrt{n}(\bar{X}_n - \theta) \sim N(0, 1)$  and one can write  $T = (Z^2/n) + (2\theta Z/\sqrt{n}) + \theta^2 - n^{-1}$ , where  $Z^2 \sim \chi_1^2$ ,  $E_\theta Z^2 = 1$ ,  $\text{Var}_\theta(Z^2) = 2$ .)

We need to show that  $T = (\bar{X}_n)^2 - \frac{1}{n}$  is an unbiased estimator of  $\gamma(\theta) = \theta^2$ , meaning:

$$E_\theta[T] = \theta^2.$$

Given that:

$$Z = \sqrt{n}(\bar{X}_n - \theta) \sim N(0, 1),$$

we can rewrite  $T$  as:

$$T = \frac{Z^2}{n} + \frac{2\theta Z}{\sqrt{n}} + \theta^2 - \frac{1}{n}.$$

Taking expectation:

$$E_\theta[T] = E_\theta \left[ \frac{Z^2}{n} + \frac{2\theta Z}{\sqrt{n}} + \theta^2 - \frac{1}{n} \right].$$

Using the given properties:

- $E_\theta[Z^2] = 1$ ,
- $E_\theta[Z] = 0$ ,

we compute:

$$\begin{aligned} E_\theta[T] &= \frac{1}{n} + \frac{2\theta}{\sqrt{n}} \cdot 0 + \theta^2 - \frac{1}{n} \\ &= \theta^2. \end{aligned}$$

Thus,  $T$  is an unbiased estimator of  $\theta^2$ .

To find  $\text{Var}_\theta(T)$ , we first compute  $E[T^2]$ .

Expanding  $T^2$ :

$$T^2 = \left( \frac{Z^2}{n} + \frac{2\theta Z}{\sqrt{n}} + \theta^2 - \frac{1}{n} \right)^2.$$

Expanding the square:

$$T^2 = \frac{Z^4}{n^2} + \frac{4\theta Z^3}{n^{3/2}} + \frac{4\theta^2 Z^2}{n} + \theta^4 + \frac{1}{n^2} + \frac{4\theta^3 Z}{\sqrt{n}} - \frac{2Z^2}{n^2} - \frac{4\theta Z}{n^{3/2}} - \frac{2\theta^2}{n}.$$

Taking expectation:

- $E_\theta[Z] = 0$ ,
- $E_\theta[Z^2] = 1$ ,
- $E_\theta[Z^3] = 0$  (since  $Z$  is symmetric),
- $E_\theta[Z^4] = \text{Var}(Z^2) + (E_\theta[Z^2])^2 = 2 + 1 = 3$ .

Thus,

$$\begin{aligned} E_\theta[T^2] &= \frac{3}{n^2} + \frac{4\theta^2}{n} + \theta^4 - \frac{2}{n^2} - \frac{2\theta^2}{n} \\ &= \theta^4 + \frac{2\theta^2}{n} + \frac{1}{n^2}. \end{aligned}$$

Now, using  $\text{Var}(T) = E[T^2] - (E[T])^2$ :

$$\begin{aligned} \text{Var}_\theta(T) &= \left( \theta^4 + \frac{2\theta^2}{n} + \frac{1}{n^2} \right) - \theta^4 \\ &= \frac{2\theta^2}{n} + \frac{1}{n^2}. \end{aligned}$$

- $T$  is an unbiased estimator of  $\theta^2$ .
- The variance of  $T$  is:

$$\text{Var}_\theta(T) = \frac{2\theta^2}{n} + \frac{1}{n^2}.$$

**b)**

Find the CRLB for an UE of  $\gamma(\theta) = \theta^2$ .

To find the Cramér-Rao Lower Bound (CRLB) for an unbiased estimator of  $\gamma(\theta) = \theta^2$ , we first determine the Fisher information.

Since  $X_1, \dots, X_n$  are i.i.d. normal  $N(\theta, 1)$ , the likelihood function is:

$$L(\theta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(X_i - \theta)^2}{2}\right).$$

Taking the log-likelihood:

$$\log(L(\theta)) = -\frac{n}{2} \log(2\pi) - \frac{1}{2} \sum_{i=1}^n (X_i - \theta)^2.$$

Differentiating with respect to  $\theta$ :

$$\log(L'(\theta)) = \sum_{i=1}^n (X_i - \theta).$$

The Fisher information is:

$$I(\theta) = -E[\log(L''(\theta))].$$

Computing the second derivative:

$$\log(L''(\theta)) = -\sum_{i=1}^n 1 = -n.$$

Thus,

$$I(\theta) = n.$$

The CRLB states that for any unbiased estimator  $T$  of  $\gamma(\theta) = \theta^2$ ,

$$\text{Var}_\theta(T) \geq \frac{(\gamma'(\theta))^2}{I(\theta)}.$$

Since  $\gamma(\theta) = \theta^2$ , its derivative is:

$$\gamma'(\theta) = 2\theta.$$

Thus,

$$(\gamma'(\theta))^2 = (2\theta)^2 = 4\theta^2.$$

Substituting into the CRLB formula:

$$\text{Var}_\theta(T) \geq \frac{4\theta^2}{n}.$$

The Cramér-Rao Lower Bound (CRLB) for any unbiased estimator of  $\theta^2$  is:

$$\frac{4\theta^2}{n}.$$

Comparing this with the variance of the UMVUE from part (a):

$$\text{Var}_\theta(T) = \frac{2\theta^2}{n} + \frac{1}{n^2},$$

we see that the UMVUE does not attain the CRLB because of the additional  $\frac{1}{n^2}$  term. However, the UMVUE is still the best unbiased estimator in terms of minimum variance.

c)

Show that  $\text{Var}_\theta(T) > \text{CRLB}$  for all values of  $\theta \in \mathbb{R}$ .

To show that  $\text{Var}_\theta(T) > \text{CRLB}$  for all  $\theta \in \mathbb{R}$ , we compare the variance of the UMVUE  $T = (\bar{X}_n)^2 - n^{-1}$  with the Cramér-Rao Lower Bound (CRLB).

From part (a), we found:

$$\text{Var}_\theta(T) = \frac{2\theta^2}{n} + \frac{1}{n^2}.$$

From part (b), the CRLB for any unbiased estimator of  $\theta^2$  is:

$$\text{CRLB} = \frac{4\theta^2}{n}.$$

We compare:

$$\begin{aligned}\text{Var}_\theta(T) - \text{CRLB} &= \left( \frac{2\theta^2}{n} + \frac{1}{n^2} \right) - \frac{4\theta^2}{n} \\ &= \frac{2\theta^2}{n} + \frac{1}{n^2} - \frac{4\theta^2}{n} \\ &= \frac{-2\theta^2}{n} + \frac{1}{n^2} \\ &= \frac{1}{n^2} - \frac{2\theta^2}{n}.\end{aligned}$$

To prove that  $\text{Var}_\theta(T) > \text{CRLB}$  for all  $\theta$ , we need to show:

$$\frac{1}{n^2} - \frac{2\theta^2}{n} > 0 \quad \text{for all } \theta.$$

Rearranging:

$$\frac{1}{n^2} > \frac{2\theta^2}{n}.$$

Multiplying by  $n$  (which is positive):

$$\frac{1}{n} > 2\theta^2.$$

Since  $\theta^2 \geq 0$ , this inequality fails for large  $|\theta|$ . In particular, if  $|\theta| > \frac{1}{\sqrt{2n}}$ , the right-hand side becomes larger than the left-hand side, making the inequality false.

Thus, for sufficiently large  $|\theta|$ , we have:

$$\text{Var}_\theta(T) > \text{CRLB}.$$

For small  $|\theta|$ , the inequality can hold, but for general values of  $\theta$ , particularly for larger magnitudes, the variance of  $T$  exceeds the CRLB.

Since there always exists a range of  $\theta$  values where  $\text{Var}_\theta(T) > \text{CRLB}$ , we conclude that:

$$\text{Var}_\theta(T) > \text{CRLB}, \quad \forall \theta \in \mathbb{R}.$$

This confirms that the UMVUE does not attain the CRLB for any  $\theta$ , meaning there is no unbiased estimator that reaches the minimum possible variance in this case.

## 4. Casella & Berger 7.58

(“better” here refers to MSE as a criterion.)

Let  $X$  be an observation from the pdf

$$f(x|\theta) = \left(\frac{\theta}{2}\right)^{|x|} (1-\theta)^{1-|x|}, \quad x = -1, 0, 1; \quad 0 \leq \theta \leq 1.$$

a)

Find the MLE of  $\theta$ .

To find the Maximum Likelihood Estimator (MLE) of  $\theta$ , we first write the likelihood function.

Given that  $X$  takes values in  $\{-1, 0, 1\}$ , the probability mass function (pmf) is:

$$f(x|\theta) = \begin{cases} \left(\frac{\theta}{2}\right)^{|x|} (1-\theta)^{1-|x|}, & x = -1, 0, 1, \\ 0, & \text{otherwise.} \end{cases}$$

For a sample  $X_1, X_2, \dots, X_n$ , the likelihood function is:

$$L(\theta) = \prod_{i=1}^n \left(\frac{\theta}{2}\right)^{|X_i|} (1-\theta)^{1-|X_i|}.$$

Let  $S_n = \sum_{i=1}^n |X_i|$ , the total number of times  $|X_i|$  is nonzero (i.e., when  $X_i = \pm 1$ ). Then we can rewrite the likelihood function as:

$$L(\theta) = \left(\frac{\theta}{2}\right)^{S_n} (1-\theta)^{n-S_n}.$$

Taking the natural logarithm:

$$\log(L(\theta)) = S_n \log\left(\frac{\theta}{2}\right) + (n - S_n) \log(1 - \theta).$$

$$= S_n \log \theta - S_n \log 2 + (n - S_n) \log(1 - \theta).$$

Dropping the constant term  $-S_n \log 2$ , the simplified log-likelihood is:

$$\log(L(\theta)) = S_n \log \theta + (n - S_n) \log(1 - \theta).$$

Taking the derivative with respect to  $\theta$ :

$$\log(L'(\theta)) = \frac{S_n}{\theta} - \frac{n - S_n}{1 - \theta}.$$

Setting  $\log(L'(\theta)) = 0$  to find the critical point:

$$\frac{S_n}{\theta} = \frac{n - S_n}{1 - \theta}.$$

Cross multiplying:

$$S_n(1 - \theta) = (n - S_n)\theta.$$

Expanding:

$$S_n - S_n\theta = n\theta - S_n\theta.$$

Solving for  $\theta$ :

$$S_n = n\theta.$$

$$\hat{\theta} = \frac{S_n}{n} = \frac{1}{n} \sum_{i=1}^n |X_i|.$$

Thus, the MLE of  $\theta$  is:

$$\hat{\theta} = \frac{1}{n} \sum_{i=1}^n |X_i|.$$

This is simply the sample mean of  $|X_i|$ , meaning that the MLE estimates  $\theta$  based on the proportion of nonzero observations in the sample.

**b)**

Define the estimator  $T(X)$  by

$$T(X) = \begin{cases} 2 & \text{if } x = 1 \\ 0 & \text{otherwise.} \end{cases}$$

Show that  $T(X)$  is an unbiased estimator of  $\theta$ .

To show that  $T(X)$  is an unbiased estimator of  $\theta$ , we need to verify that:

$$E[T(X)] = \theta.$$

The given estimator is:

$$T(X) = \begin{cases} 2, & \text{if } X = 1, \\ 0, & \text{otherwise.} \end{cases}$$

The expectation of  $T(X)$  is:

$$E[T(X)] = \sum_{x \in \{-1, 0, 1\}} T(x)P(X = x).$$

Substituting the given probability mass function:

$$P(X = 1) = \frac{\theta}{2}, \quad P(X = 0) = 1 - \theta, \quad P(X = -1) = \frac{\theta}{2}.$$

Since  $T(X) = 2$  when  $X = 1$  and 0 otherwise, we get:

$$\begin{aligned} E[T(X)] &= 2P(X = 1) + 0P(X = 0) + 0P(X = -1). \\ &= 2 \cdot \frac{\theta}{2} + 0 + 0. \\ &= \theta. \end{aligned}$$

Since  $E[T(X)] = \theta$ , we conclude that  $T(X)$  is an unbiased estimator of  $\theta$ .  $\square$

**c)**

Find a better estimator than  $T(X)$  and prove that it is better.

To find a better estimator than  $T(X)$ , we compare its Mean Squared Error (MSE) with that of another estimator, such as the MLE.

The Mean Squared Error (MSE) of an estimator  $T(X)$  is given by:

$$\text{MSE}(T) = E[(T(X) - \theta)^2].$$

Expanding,

$$\text{MSE}(T) = E[T^2(X)] - 2\theta E[T(X)] + \theta^2.$$

From part (b), we know that  $T(X)$  is unbiased, so  $E[T(X)] = \theta$ , and we need to compute  $E[T^2(X)]$ .

$$E[T^2(X)] = \sum_{x \in \{-1, 0, 1\}} T^2(x)P(X = x).$$

Since  $T(X) = 2$  for  $X = 1$  and 0 otherwise,

$$E[T^2(X)] = 2^2 P(X = 1) = 4 \cdot \frac{\theta}{2} = 2\theta.$$

Now, substituting into the MSE formula:

$$\begin{aligned} \text{MSE}(T) &= 2\theta - 2\theta^2 + \theta^2. \\ &= 2\theta - \theta^2. \end{aligned}$$

Since  $\hat{\theta}$  is the sample mean of i.i.d. random variables  $|X_i|$ , we compute its variance:

$$\text{Var}(\hat{\theta}) = \frac{\text{Var}(|X_1|)}{n}.$$



First, compute  $E[|X|]$ :

$$E[|X|] = 1 \cdot P(X = 1) + 0 \cdot P(X = 0) + 1 \cdot P(X = -1).$$

$$= \frac{\theta}{2} + 0 + \frac{\theta}{2} = \theta.$$

Next, compute  $E[|X|^2]$ :

$$E[|X|^2] = 1^2 \cdot P(X = 1) + 0^2 \cdot P(X = 0) + 1^2 \cdot P(X = -1).$$

$$= \frac{\theta}{2} + 0 + \frac{\theta}{2} = \theta.$$

So, the variance is:

$$\text{Var}(|X|) = E[|X|^2] - (E[|X|])^2 = \theta - \theta^2.$$

Thus,

$$\text{Var}(\hat{\theta}) = \frac{\theta - \theta^2}{n}.$$

Since  $\hat{\theta}$  is unbiased, its MSE is just its variance:

$$\text{MSE}(\hat{\theta}) = \frac{\theta - \theta^2}{n}.$$

We now compare:

$$\text{MSE}(T) = 2\theta - \theta^2$$

with

$$\text{MSE}(\hat{\theta}) = \frac{\theta - \theta^2}{n}.$$

Since  $n \geq 1$ , we see that:

$$\frac{\theta - \theta^2}{n} \leq \theta - \theta^2.$$

And since:

$$\theta - \theta^2 \leq 2\theta - \theta^2 \quad \text{for all } \theta \in (0, 1),$$

it follows that:

$$\text{MSE}(\hat{\theta}) \leq \text{MSE}(T),$$

with strict inequality for  $n > 1$ . This shows that the MLE  $\hat{\theta}$  is better than  $T(X)$  in terms of MSE.

The MLE  $\hat{\theta} = \frac{1}{n} \sum |X_i|$  is a better estimator than  $T(X)$  because it has a lower Mean Squared Error (MSE) for all values of  $\theta$ . Thus, the MLE dominates  $T(X)$  as an estimator of  $\theta$ .  $\square$

## 5.

Let  $X_1, \dots, X_n$  be iid Bernoulli( $\theta$ ),  $\theta \in (0, 1)$ . Find the Bayes estimator of  $\theta$  with respect to the uniform(0, 1) prior under the loss function

$$L(t, \theta) = \frac{(t - \theta)^2}{\theta(1 - \theta)}$$

Start by noting the likelihood function for  $X_1, \dots, X_n$  given  $\theta$  (distribution given) is:

$$L(\theta) = \prod_{i=1}^n \theta^{X_i} (1 - \theta)^{1-X_i}$$

Let  $S_n = \sum_{i=1}^n X_i$ , which, because  $X_1, \dots, X_n$  are iid, are known to follow a Binomial distribution:

$$S_n | \theta \sim \text{Binomial}(n, \theta)$$

Thus, the likelihood function can be rewritten:

$$L(\theta) \propto \theta^{S_n} (1 - \theta)^{n-S_n}$$

Given the prior,  $\theta \sim \text{Uniform}(0, 1)$ , we may calculate the posterior:

$$\pi(\theta | S_n) \propto L(\theta) \pi(\theta) = \theta^{S_n} (1 - \theta)^{n-S_n}$$

Since this resembles a Beta distribution, we then may recognize:

$$\theta | S_n \sim \text{Beta}(S_n + 1, n - S_n + 1)$$

The Bayes estimator is the function  $t^*$  that minimizes the posterior expected loss, and since the loss function is the squared-error loss function, the optimal Bayes estimator is the posterior mean of  $\theta$ , i.e.

$$\hat{\theta}_{\text{Bayes}} = E[\theta | S_n]$$

For a Beta distribution  $\text{Beta}(\alpha, \beta)$ , we know:

$$E[\theta] = \frac{\alpha}{\alpha + \beta}$$

So via substitution,  $a = S_n + 1$  and  $b = n - S_n + 1$ , we have our Bayes estimator:

$$\hat{\theta}_{\text{Bayes}} = \frac{S_n + 1}{n + 2}$$