

HW3

Sam Olson

Outline

- Q1: g2g
- Q2: g2g
- Q3: part c) ending a bit wonky
- Q4: g2g
- Q5: wonky with \propto usage

1.

Suppose X_1, \dots, X_n are iid Bernoulli(p), $0 < p < 1$.

a)

Find the information number $I_n(p)$ and make a rough sketch of $I_n(p)$ as a function of $p \in (0, 1)$.

Given that X_1, \dots, X_n are i.i.d. Bernoulli(p), the likelihood function is:

$$L(p) = \prod_{i=1}^n p^{X_i} (1-p)^{1-X_i}$$

Taking the log-likelihood,

$$\log(L(p)) = \sum_{i=1}^n [X_i \log p + (1 - X_i) \log(1 - p)]$$

The first derivative is:

$$\log(L(p))' = \sum_{i=1}^n \left[\frac{X_i}{p} - \frac{1 - X_i}{1 - p} \right] = \sum_{i=1}^n \frac{X_i - p}{p(1 - p)}$$

The Fisher information is:

$$I_n(p) = -E[\log(L(p))'']$$

Computing the second derivative:

$$\log(L(p))'' = \sum_{i=1}^n \left[-\frac{X_i}{p^2} - \frac{1 - X_i}{(1 - p)^2} \right]$$

Taking expectation:

$$E[\log(L(p))''] = \sum_{i=1}^n \left[-\frac{E[X_i]}{p^2} - \frac{E[1 - X_i]}{(1-p)^2} \right]$$

Given we know the distribution of the random variables, we know $E[X_i] = p$ and $E[1 - X_i] = 1 - p$. This allows us to simplify the expression:

$$E[\log(L(p))''] = \sum_{i=1}^n \left[-\frac{p}{p^2} - \frac{1-p}{(1-p)^2} \right] = \sum_{i=1}^n \left[-\frac{1}{p} - \frac{1}{1-p} \right] = -n \left[\frac{1}{p} + \frac{1}{1-p} \right]$$

Noting linearity of Fisher information:

$$I_n(p) = n \left[\frac{1}{p} + \frac{1}{1-p} \right]$$


Sketch

```
# functional form
fisher_info <- function(p, n) {
  return(n * (1/p + 1/(1 - p)))
}

# setup
p_values <- seq(0.01, 0.99, length.out = 100)
n <- 1
I_values <- fisher_info(p_values, n)

# plot
plot(x = p_values,
     y = I_values,
     type = "l",
     col = "black", lwd = 2,
     xlab = "p", ylab = "Fisher Information",
     main = "Fisher Information for Bernoulli(p)")
abline(v = 0.5, lty = 2, col = "gray")
```

Fisher Information for Bernoulli(p)



b)

Find the value of $p \in (0, 1)$ for which $I_n(p)$ is minimal. (This value of p corresponds to the “hardest” case for estimating p . That is, when data are generated under this value of p from the model, the variance of an UE of p is potentially largest.)

To find the value of p that minimizes the Fisher information $I_n(p)$, we use the functional form of the Fisher Information:

$$I_n(p) = n \left[\frac{1}{p} + \frac{1}{1-p} \right]$$

Differentiating $I_n(p)$ with respect to p , and setting equal to zero:

$$I_n(p)' = n \left[-\frac{1}{p^2} + \frac{1}{(1-p)^2} \right] = -\frac{1}{p^2} + \frac{1}{(1-p)^2} = 0$$

This gives us the expression:

$$\frac{1}{p^2} = \frac{1}{(1-p)^2}$$

Taking square roots:

$$\frac{1}{p} = \frac{1}{1-p} \rightarrow p = 1-p \rightarrow p = \frac{1}{2}$$

To ensure this is a maximum, we also check whether the second derivative is positive (since we are minimizing and not maximizing) at $\frac{1}{2}$:

$$I_n(p)' = n \left[\frac{2}{p^3} + \frac{2}{(1-p)^3} \right]$$

$$I_n \left(\frac{1}{2} \right)'' = n \left[\frac{2}{(1/2)^3} + \frac{2}{(1/2)^3} \right] = n \left[\frac{2}{1/8} + \frac{2}{1/8} \right] = n [16 + 16] = 32n > 0$$

So this is in fact a minimum, hence the Fisher information is minimized at:

$$p = \frac{1}{2}$$

c)

Show that $\hat{X}_n = \sum_{i=1}^n X_i/n$ is the UMVUE of p .

Note to self: Uniformly Minimum Variance Unbiased Estimator (UMVUE)

We start by checking if \hat{X}_n is an unbiased estimator of p :

$$E[\hat{X}_n] = E \left[\frac{1}{n} \sum_{i=1}^n X_i \right] = \frac{1}{n} \sum_{i=1}^n E[X_i] E[\hat{X}_n] = \frac{1}{n} \cdot np = p$$

$$Bias(\bar{X}_n) = E[\hat{X}_n] - E[X] = p - p = 0$$

So \hat{X}_n is an unbiased estimator of p .

Now as far as the “Uniformly Minimum Variance” part of the question:

Note again the Fisher Information formula we’ve found:

$$I_n(p) = \frac{np}{p^2} + \frac{n(1-p)}{(1-p)^2} = \frac{n}{p(1-p)}$$

By the definition, the Cramér-Rao Lower Bound, for any unbiased estimator T of p :

$$\text{Var}_p(T) \geq \frac{(\gamma'(p))^2}{I_n(p)}$$

Here, we are estimating $\gamma(p) = p$, so $\gamma'(p) = 1$. Therefore:

$$\text{Var}_p(T) \geq \frac{1^2}{I_n(p)} = \frac{p(1-p)}{n}$$

We compute the variance of $\hat{X}_n = S_n/n$:

$$E[\hat{X}_n] = E \left[\frac{S_n}{n} \right] = \frac{1}{n} E[S_n] = \frac{np}{n} = p$$

$$\text{Var}(\hat{X}_n) = \text{Var} \left(\frac{S_n}{n} \right) = \frac{1}{n^2} \text{Var}(S_n)$$

Since $S_n \sim \text{Binomial}(n, p)$, we know:

$$\text{Var}(S_n) = np(1 - p)$$

Thus:

$$\text{Var}(\hat{X}_n) = \frac{np(1 - p)}{n^2} = \frac{p(1 - p)}{n}$$

Comparing with the CRLB:

$$\text{Var}(\hat{X}_n) = \frac{p(1 - p)}{n} = \frac{1}{I_n(p)}$$

Since \hat{X}_n attains the bound, it is an efficient estimator.

Since \hat{X}_n is unbiased and attains the CRLB, it is the UMVUE.

2.

Suppose that the random variables Y_1, \dots, Y_n satisfy

$$Y_i = \beta x_i + \varepsilon_i, \quad i = 1, \dots, n$$

where x_1, \dots, x_n are fixed constants and $\varepsilon_1, \dots, \varepsilon_n$ are iid $N(0, \sigma^2)$; here we assume $\sigma^2 > 0$ is known.

a)

Find the MLE of β .

To find the Maximum Likelihood Estimator (MLE) of β , we first write the likelihood function.

Since $Y_i = \beta x_i + \varepsilon_i$ with $\varepsilon_i \sim N(0, \sigma^2)$, we have:

$$Y_i \sim N(\beta x_i, \sigma^2)$$

Thus, the joint density function of Y_1, \dots, Y_n is:

$$L(\beta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(Y_i - \beta x_i)^2}{2\sigma^2}\right)$$

Taking the log-likelihood:

$$\log(L(\beta)) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (Y_i - \beta x_i)^2$$

To find the MLE of β , we take the derivative with respect to β and set to zero:

$$\frac{d}{d\beta} \log(L(\beta)) = -\frac{1}{2\sigma^2} \cdot (-2) \sum_{i=1}^n x_i (Y_i - \beta x_i) = \frac{1}{\sigma^2} \sum_{i=1}^n x_i (Y_i - \beta x_i) \rightarrow \sum_{i=1}^n x_i Y_i - \beta \sum_{i=1}^n x_i^2 = 0$$

Solving for β , we get our MLE of β as::

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i Y_i}{\sum_{i=1}^n x_i^2}$$

To ensure this is a maximum, we take the second derivative and see if it is negative:

$$\log(L(\beta))'' = -\sum_{i=1}^n x_i^2 < 0$$

So this is in fact the maximum.

b)

Find the distribution of the MLE.

From part a), the MLE of β is:

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i Y_i}{\sum_{i=1}^n x_i^2}$$

To determine the distribution of $\hat{\beta}$, determine its expectation and variance, noting that since $\hat{\beta}$ is a linear combination of the normal random variables ε_i , it follows that $\hat{\beta}$ itself is normally distributed.

That being said, given $Y_i = \beta x_i + \varepsilon_i$, we may write:

$$\hat{\beta} = \frac{\sum_{i=1}^n x_i(\beta x_i + \varepsilon_i)}{\sum_{i=1}^n x_i^2} = \frac{\beta \sum_{i=1}^n x_i^2 + \sum_{i=1}^n x_i \varepsilon_i}{\sum_{i=1}^n x_i^2}$$

Taking the expectation, noting our data is treated as “fixed”, we may write:

$$E[\hat{\beta}] = \frac{\beta \sum_{i=1}^n x_i^2 + \sum_{i=1}^n x_i E[\varepsilon_i]}{\sum_{i=1}^n x_i^2} = \frac{\beta \sum_{i=1}^n x_i^2}{\sum_{i=1}^n x_i^2} = \beta$$

Noting $E[\varepsilon_i] = 0$

Because $E[\hat{\beta}] = \beta$, it has zero bias and $\hat{\beta}$ is an unbiased estimator of β . Not needed for the distribution, but will need this note for later.

Let us then analyze the variance. We start again with definitions:

$$\text{Var}(\hat{\beta}) = \text{Var}\left(\beta + \frac{\sum_{i=1}^n x_i \varepsilon_i}{\sum_{i=1}^n x_i^2}\right) = \text{Var}(\beta) + \text{Var}\left(\frac{\sum_{i=1}^n x_i \varepsilon_i}{\sum_{i=1}^n x_i^2}\right)$$

Simplifying:

$$\text{Var}(\hat{\beta}) = \text{Var}\left(\frac{\sum_{i=1}^n x_i \varepsilon_i}{\sum_{i=1}^n x_i^2}\right) = \text{Var}\left(\sum_{i=1}^n \frac{(x_i^2 \sigma^2)}{(x_i^2)^2}\right) = \frac{\sigma^2}{\sum_{i=1}^n x_i^2}$$

We thus conclude:

$$\hat{\beta} \sim N\left(\beta, \frac{\sigma^2}{\sum_{i=1}^n x_i^2}\right)$$

c)

Find the CRLB for estimating β . (Hint: you'll have to work with the joint distribution $f(y_1, \dots, y_n | \beta)$ directly, since Y_1, \dots, Y_n are not iid.)

To find the CRLB, we first calculate the Fisher information.

Note the joint density:

$$f(Y_1, \dots, Y_n | \beta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(Y_i - \beta x_i)^2}{2\sigma^2}\right)$$

Taking the log-likelihood:

$$\log(L(\beta)) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (Y_i - \beta x_i)^2$$

We take the derivative:

$$\log(L(\beta))' = -\frac{1}{2\sigma^2} \cdot (-2) \sum_{i=1}^n x_i (Y_i - \beta x_i) = \frac{1}{\sigma^2} \sum_{i=1}^n x_i (Y_i - \beta x_i)$$

The Fisher information is then:

$$I(\beta) = -E[\log(L(\beta))''] = -E \left[-\frac{1}{\sigma^2} \sum_{i=1}^n x_i^2 \right] = \frac{1}{\sigma^2} \sum_{i=1}^n x_i^2$$

We then have what we need to calculate the CRLB using the information we've gathered.

The CRLB is:

$$\frac{1}{I(\beta)} = \frac{1}{\frac{1}{\sigma^2} \sum_{i=1}^n x_i^2} = \frac{\sigma^2}{\sum_{i=1}^n x_i^2}$$

d)

Show the MLE is the UMVUE of β .

Now we just need to compare the variance of our MLE of β to the value calculated in part c). To that end:

We have already calculated the expectation of $\hat{\beta}_{MLE}$, which is β , so via Bias calculation:

$$\text{Bias}(\hat{\beta}_{MLE}) = E[\hat{\beta}_{MLE}] - \beta = \beta - \beta = 0$$

Hence it is unbiased. We then just need to determine if our MLE attains the CRLB. If so, then the MLE is the UMVUE.

Recall the variance of the MLE:

$$\frac{\sigma^2}{\sum_{i=1}^n x_i^2}$$

And the CRLB:

$$\frac{\sigma^2}{\sum_{i=1}^n x_i^2}$$

These are one and the same! So we do indeed satisfy:

$$\text{Var}(\hat{\beta}_{MLE}) = CRLB$$

Such that the MLE is the UMVUE.

3.

Suppose X_1, \dots, X_n are iid normal $N(0, 1)$, where $\theta \in \mathbb{R}$. It turns out that $T = (\bar{X}_n)^2 - n^{-1}$ is the UMVUE of $\gamma(\theta) = \theta^2$. (We can show this later in the course; our goal here is to show that the UMVUE can exist without obtaining the CRLB.)

a)

Show T is an UE of $\gamma(\theta) = \theta^2$ and find the variance $\text{Var}_\theta(T)$ of T . (Note $Z = \sqrt{n}(\bar{X}_n - \theta) \sim N(0, 1)$ and one can write $T = (Z^2/n) + (2\theta Z/\sqrt{n}) + \theta^2 - n^{-1}$, where $Z^2 \sim \chi_1^2$, $E_\theta Z^2 = 1$, $\text{Var}_\theta(Z^2) = 2$.)

Given:

$$Z = \sqrt{n}(\bar{X}_n - \theta) \sim N(0, 1)$$

we can rewrite T in terms of Z , specifically:

$$T = \frac{Z^2}{n} + \frac{2\theta Z}{\sqrt{n}} + \theta^2 - \frac{1}{n}$$

Taking expectation:

$$E_\theta[T] = E_\theta \left[\frac{Z^2}{n} + \frac{2\theta Z}{\sqrt{n}} + \theta^2 - \frac{1}{n} \right] = \frac{1}{n} + \frac{2\theta}{\sqrt{n}}(0) + \theta^2 - \frac{1}{n} = \theta^2$$

Thus, T is an unbiased estimator of θ^2 .

We then must calculate the variance of T , to that end, we find $E[T^2]$:

As defined:

$$T^2 = \left(\frac{Z^2}{n} + \frac{2\theta Z}{\sqrt{n}} + \theta^2 - \frac{1}{n} \right)^2 = \frac{Z^4}{n^2} + \frac{4\theta Z^3}{n^{3/2}} + \frac{4\theta^2 Z^2}{n} + \theta^4 + \frac{1}{n^2} + \frac{4\theta^3 Z}{\sqrt{n}} - \frac{2Z^2}{n^2} - \frac{4\theta Z}{n^{3/2}} - \frac{2\theta^2}{n}$$

Though that's quite a lot, we can actually simplify it quite a bit when taking expectation, noting the distribution of Z aids in these calculations

(Note: $E_\theta[Z] = 0$, $E_\theta[Z^2] = 1$, $E_\theta[Z^3] = 0$, and $E_\theta[Z^4] = \text{Var}(Z^2) + (E_\theta[Z^2])^2 = 2 + 1 = 3$.)

Thus,

$$E_\theta[T^2] = \frac{3}{n^2} + \frac{4\theta^2}{n} + \theta^4 - \frac{2}{n^2} - \frac{2\theta^2}{n} = \theta^4 + \frac{2\theta^2}{n} + \frac{1}{n^2}$$

Now we can calculate the variance:

$$\text{Var}_\theta(T) = \left(\theta^4 + \frac{2\theta^2}{n} + \frac{1}{n^2} \right) - \theta^4 = \frac{2\theta^2}{n} + \frac{1}{n^2}$$

b)

Find the CRLB for an UE of $\gamma(\theta) = \theta^2$.

Since X_1, \dots, X_n are i.i.d. normal $N(\theta, 1)$, the likelihood function is:

$$L(\theta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{(X_i - \theta)^2}{2}\right)$$

Taking the log-likelihood:

$$\log(L(\theta)) = -\frac{n}{2} \log(2\pi) - \frac{1}{2} \sum_{i=1}^n (X_i - \theta)^2$$

Differentiating with respect to θ , and getting expectation to derive the Fisher Information:

$$\log(L(\theta))' = \sum_{i=1}^n (X_i - \theta) \rightarrow I(\theta) = -E[\log(L(\theta))'] = -E\left[-\sum_{i=1}^n 1\right] = -(-n) = n$$

The CRLB by definition is given by:

$$\frac{(\gamma'(\theta))^2}{I(\theta)}$$

We just need now to calculate the numerator. To that end, note that $\gamma(\theta) = \theta^2$, making its derivative:

$$\gamma'(\theta) = 2\theta$$

Thus the CRLB is:

$$\frac{(\gamma'(\theta))^2}{n} = \frac{(2\theta)^2}{n} = \frac{4\theta^2}{n}$$

c)

Show that $\text{Var}_\theta(T) > \text{CRLB}$ for all values of $\theta \in \mathbb{R}$.

We are now tasked with comparing the variance of the UMVUE $T = (\bar{X}_n)^2 - n^{-1}$ with the Cramér-Rao Lower Bound (CRLB) from part b). To that end note our prior results:

From part a):

$$\text{Var}_\theta(T) = \frac{2\theta^2}{n} + \frac{1}{n^2}$$

From part b), the CRLB (for any unbiased estimator of θ^2) is:

$$\text{CRLB} = \frac{4\theta^2}{n}$$

Comparing these two quantities directly, their difference is given by:

$$\text{Var}_\theta(T) - \text{CRLB} = \left(\frac{2\theta^2}{n} + \frac{1}{n^2} \right) - \frac{4\theta^2}{n} = \frac{2\theta^2}{n} + \frac{1}{n^2} - \frac{4\theta^2}{n} = \frac{-2\theta^2}{n} + \frac{1}{n^2} = \frac{1}{n^2} - \frac{2\theta^2}{n}.$$

To establish whether $\text{Var}_\theta(T) > \text{CRLB}$ for all θ , we analyze the sign of:

$$\frac{1}{n^2} - \frac{2\theta^2}{n}.$$

Since $n > 0$, we can rewrite the inequality as:

$$\frac{1}{n^2} > \frac{2\theta^2}{n} \quad \Rightarrow \quad \frac{1}{n} > 2\theta^2.$$

Since $\theta^2 \geq 0$, this inequality only holds for sufficiently small $|\theta|$. In particular, if:

$$|\theta| > \frac{1}{\sqrt{2n}},$$

then the right-hand side of the expression exceeds the left-hand side, making the inequality false.

Thus, for sufficiently large $|\theta|$, we have:

$$\text{Var}_\theta(T) > \text{CRLB}.$$

For small $|\theta|$, equality can hold, but in general, for larger magnitudes of θ , the variance of T exceeds the CRLB.

Since there always exists a range of θ values where $\text{Var}_\theta(T) > \text{CRLB}$, we conclude that:

$$\text{Var}_\theta(T) \geq \text{CRLB}, \quad \forall \theta \in \mathbb{R},$$

with strict inequality for sufficiently large $|\theta|$.

4. Casella & Berger 7.58

(“better” here refers to MSE as a criterion.)

Let X be an observation from the pdf

$$f(x|\theta) = \left(\frac{\theta}{2}\right)^{|x|} (1-\theta)^{1-|x|}, \quad x = -1, 0, 1; \quad 0 \leq \theta \leq 1.$$

a)

Find the MLE of θ .

Given that X takes values in $\{-1, 0, 1\}$, it is discrete, so we note the pmf:

$$f(x|\theta) = \begin{cases} \left(\frac{\theta}{2}\right)^{|x|} (1-\theta)^{1-|x|} & x = -1, 0, 1 \\ 0 & \text{otherwise} \end{cases}$$

For a sample X_1, X_2, \dots, X_n , the likelihood function is:

$$L(\theta) = \prod_{i=1}^n \left(\frac{\theta}{2}\right)^{|X_i|} (1-\theta)^{1-|X_i|}$$

Let $S_n = \sum_{i=1}^n |X_i|$. We may then rewrite the likelihood function as:

$$L(\theta) = \left(\frac{\theta}{2}\right)^{S_n} (1-\theta)^{n-S_n}$$

Using our log-likelihood technique:

$$\log(L(\theta)) = S_n \log\left(\frac{\theta}{2}\right) + (n-S_n) \log(1-\theta) = S_n \log \theta - S_n \log 2 + (n-S_n) \log(1-\theta) = S_n \log \theta + (n-S_n) \log(1-\theta)$$

We find the maximum the typical route, i.e., taking the derivative with respect to θ and setting equal to zero:

$$\log(L(\theta))' = \frac{S_n}{\theta} - \frac{n-S_n}{1-\theta} = 0 \rightarrow \frac{S_n}{\theta} = \frac{n-S_n}{1-\theta}$$

After some simplifying:

$$S_n(1-\theta) = (n-S_n)\theta \rightarrow S_n - S_n\theta = n\theta - S_n\theta \rightarrow S_n = n\theta$$

And we arrive at our “MLE” (in quotes because there’s our second check to account for):

$$\hat{\theta} = \frac{S_n}{n} = \frac{1}{n} \sum_{i=1}^n |X_i|$$

To double check, we take the second derivative (at the MLE) and see if it is negative:

$$\log(L(\theta))'' = -\frac{S_n}{\theta^2} - \frac{n - S_n}{(1 - \theta)^2}$$

$$\log(L(\hat{\theta}))'' = -\frac{S_n n^2}{S_n^2} - \frac{(n - S_n)n^2}{(n - S_n)^2} = -\frac{n^2}{S_n} - \frac{n^2}{n - S_n} < 0$$

Noting: $S_n > 0$ and $n - S_n > 0$

So yes, this is our maximum and our MLE!

Jetzt zock' ich Fortnite und trink' Cola! Yippee!

b)

Define the estimator $T(X)$ by

$$T(X) = \begin{cases} 2 & \text{if } x = 1 \\ 0 & \text{otherwise.} \end{cases}$$

Show that $T(X)$ is an unbiased estimator of θ .

To test for bias, we find the the expectation of $T(X)$:

$$E[T(X)] = \sum_{x \in \{-1, 0, 1\}} T(x)P(X = x).$$

Using the pmf from part a), each possible outcome/observation of X has its associated probability given by:

$$P(X = 1) = \frac{\theta}{2}, \quad P(X = 0) = 1 - \theta, \quad P(X = -1) = \frac{\theta}{2}$$

So we need to do the more “manual” calculation of expectation:

$$E[T(X)] = 2P(X = 1) + 0P(X = 0) + 0P(X = -1).$$

Since $T(X) = 2$ when $X = 1$ and 0 otherwise from the initial definition of T .

Thus, we calculate:

$$E[T(X)] = 2 \cdot \frac{\theta}{2} + 0 + 0 = \theta$$

So, via Bias calculation, we know $T(X)$ is an unbiased estimator of θ because $E[T(X)] = \theta$.

c)

Find a better estimator than $T(X)$ and prove that it is better.

By “better” we are making note of the “hint” to compare MSE, and “better” corresponding to smaller MSE compared to $T(X)$.

By definition, the MSE of the estimator $T(X)$ is:

$$\text{MSE}(T) = E[(T(X) - \theta)^2] = E[T^2(X)] - 2\theta E[T(X)] + \theta^2$$

From part b), we know that $T(X)$ is unbiased, so the unknown quantity in the above expression is $E[T^2(X)]$. Solving for that:

$$E[T^2(X)] = \sum_{x \in \{-1, 0, 1\}} T^2(x)P(X = x) = 2^2 P(X = 1) = 4 \cdot \frac{\theta}{2} = 2\theta$$

Returning to the MSE, our goal is to then find a better (smaller) MSE than $T(X)$, which is:

$$\text{MSE}(T) = 2\theta - 2\theta^2 + \theta^2 = \theta - \theta^2 = \theta(1 - \theta)$$

Our first guess will be to use the sample mean, the MLE from part a):

$$\hat{\theta} = \frac{S_n}{n} = \frac{1}{n} \sum_{i=1}^n |X_i|$$

To find the relevant quantities to find its MSE, we start with:

$E[|X|]$:

$$E[|X|] = 1 \cdot P(X = 1) + 0 \cdot P(X = 0) + 1 \cdot P(X = -1) = \frac{\theta}{2} + 0 + \frac{\theta}{2} = \theta$$

Next, $E[|X|^2]$:

$$E[|X|^2] = 1^2 \cdot P(X = 1) + 0^2 \cdot P(X = 0) + 1^2 \cdot P(X = -1) = \frac{\theta}{2} + 0 + \frac{\theta}{2} = \theta$$

So, the variance is:

$$\text{Var}(|X|) = E[|X|^2] - (E[|X|])^2 = \theta - \theta^2 \rightarrow \text{Var}(\hat{\theta}) = \frac{\theta - \theta^2}{n}$$

Since $\hat{\theta}$ is unbiased, i.e.

$$E[\hat{\theta}] = E\left[\frac{S_n}{n}\right] = E\left[\frac{1}{n} \sum_{i=1}^n |X_i|\right] = \frac{1}{n} E\left[\sum_{i=1}^n |X_i|\right] = \frac{n\theta}{n} = \theta$$

The MSE of $\hat{\theta}$ is:

$$\text{MSE}(\hat{\theta}) = \frac{\theta - \theta^2}{n}$$

We now comparing the two estimators:

$$\text{MSE}(T) = 2\theta - \theta^2$$

$$\text{MSE}(\hat{\theta}) = \frac{\theta - \theta^2}{n}$$

Since $n \geq 1$:

$$\frac{\theta - \theta^2}{n} \leq \theta - \theta^2$$

And since:

$$\theta - \theta^2 \leq 2\theta - \theta^2 \quad \forall \theta \in (0, 1)$$

it follows:

$$\text{MSE}(\hat{\theta}) \leq \text{MSE}(T)$$

with strict inequality for $n > 1$. So the MLE $\hat{\theta} = \frac{1}{n} \sum |X_i|$ is a “better” estimator than $T(X)$ because it has a lower Mean Squared Error for all values of θ (while also being unbiased!)

5.

Let X_1, \dots, X_n be iid Bernoulli(θ), $\theta \in (0, 1)$. Find the Bayes estimator of θ with respect to the uniform(0, 1) prior under the loss function

$$L(t, \theta) = \frac{(t - \theta)^2}{\theta(1 - \theta)}$$

Start by noting the likelihood function for X_1, \dots, X_n given θ (distribution given) is:

$$L(\theta) = \prod_{i=1}^n \theta^{X_i} (1 - \theta)^{1-X_i}$$

Let $S_n = \sum_{i=1}^n X_i$, which, because X_1, \dots, X_n are iid, are known to follow a Binomial distribution:

$$S_n | \theta \sim \text{Binomial}(n, \theta)$$

Thus, the likelihood function can be rewritten:

$$L(\theta) \propto \theta^{S_n} (1 - \theta)^{n-S_n}$$

Given the prior, $\theta \sim \text{Uniform}(0, 1)$, we may calculate the posterior:

$$\pi(\theta | S_n) \propto L(\theta) \pi(\theta) = \theta^{S_n} (1 - \theta)^{n-S_n}$$

Since this resembles a Beta distribution, we then may recognize:

$$\theta | S_n \sim \text{Beta}(S_n + 1, n - S_n + 1)$$

The Bayes estimator is the function t^* that minimizes the posterior expected loss, and since the loss function is the squared-error loss function, the optimal Bayes estimator is the posterior mean of θ , i.e.

$$\hat{\theta}_{\text{Bayes}} = E[\theta | S_n]$$

For a Beta distribution $\text{Beta}(\alpha, \beta)$, we know:

$$E[\theta] = \frac{\alpha}{\alpha + \beta}$$

So via substitution, $a = S_n + 1$ and $b = n - S_n + 1$, we have our Bayes estimator:

$$\hat{\theta}_{\text{Bayes}} = \frac{S_n + 1}{n + 2}$$