HW7

Sam Olson

Q1

Problem 8.6 a) - b), Casella and Berger (2nd Edition)

Suppose that we have two independent random samples: $X_1, \ldots, X_n$ are exponential(θ), and $Y_1, \ldots, Y_m$ are exponential(μ).

a)

Find the LRT of

$$H_0 : \theta = \mu \quad \text{versus} \quad H_1 : \theta \neq \mu.$$

Answer

The LRT statistic is of the form:

$$\lambda(x, y) = \frac{\max_\theta L(\theta \mid x, y)}{\max_{\theta,\mu} L(\theta, \mu \mid x, y)}$$

Where, under H_0 ($\theta = \mu$).

Generally, we know that, the MLE will be some weighted average of the observations, taking advantage of the one parameter exponential families known to be complete and their MLEs of a general form. That is, the MLE will be some weighted average of the observations, taking advantage of the one parameter exponential families known to be complete and their MLEs of a general form.

Under H_0 (to get the numerator of the LRT) the MLE for θ is of the form:

$$\hat{\theta}_{H_0} = \frac{\sum_{i=1}^n X_i + \sum_{j=1}^m Y_j}{n + m}$$

And, under the full model (the denominator of the LRT), the MLEs are the individual sample means, i.e.:

$$\hat{\theta} = \bar{X} = \frac{\sum X_i}{n}, \quad \hat{\mu} = \bar{Y} = \frac{\sum Y_j}{m}$$

Returning to the original LRT expression, we then have:

$$\lambda(x, y) = \frac{(\hat{\theta}_0)^{-(n+m)} e^{-(n+m)}}{(\hat{\theta})^{-n} e^{-n} (\hat{\mu})^{-m} e^{-m}} = \frac{(\bar{X})^n (\bar{Y})^m}{\left(\frac{\sum X_i + \sum Y_j}{n+m}\right)^{n+m}} = \frac{(n+m)^{n+m} (\sum X_i)^n (\sum Y_j)^m}{n^n m^m (\sum X_i + \sum Y_j)^{n+m}}$$

We may then construct our test function, where our rejection rule is to "Reject H_0 if $\lambda(x, y) \leq c$", where c is calibrated based on the significance level α, i.e.:

$$\varphi(x, y) = \begin{cases} 1 & \text{if } \lambda(x, y) \leq c, \\ 0 & \text{otherwise} \end{cases}$$

Where (to save space above), we have:

$$\lambda(x, y) = \frac{(n+m)^{n+m}(\sum X_i)^n(\sum Y_j)^m}{n^n m^m(\sum X_i + \sum Y_j)^{n+m}}$$

And c is chosen such that $P(\varphi(X, Y) = 1 \mid H_0) = \alpha$.

b)

Show that the test in part a) can be based on the statistic

$$T = \frac{\sum X_i}{\sum X_i + \sum Y_i}$$

Answer

Let $T = \frac{\sum X_i}{\sum X_i + \sum Y_j}$.

Rewriting the LRT from part a) in terms of T:

$$\lambda(x, y) = \frac{(n+m)^{n+m}}{n^n m^m}\left(\frac{\sum X_i}{\sum X_i + \sum Y_j}\right)^n\left(\frac{\sum Y_j}{\sum X_i + \sum Y_j}\right)^m = \frac{(n+m)^{n+m}}{n^n m^m}T^n(1-T)^m$$

Since $\lambda(x, y)$ depends on the data only through T, the LRT can be based entirely on T.

Using the above, we may define the rejection region where the test rejects H_0 when T is too small or too large with constants a and b, where:

$$T \leq a \quad \text{or} \quad T \geq b$$

And where a and b are critical values satisfying:

$$P(T \leq a \mid H_0) + P(T \geq b \mid H_0) = \alpha$$

Under H_0 ($\theta = \mu$), $\sum X_i \sim \text{Gamma}(n, \theta)$, $\sum Y_j \sim \text{Gamma}(m, \theta)$, as the sum of iid Exponentials is Gamma, and a linear combination of Gamma distributions with common rate parameter θ is a Beta. Since both X and Y are independent of one another, their sums are also independent, and determining the parameters of the T Beta distribution becomes a matter of algebra (the distribution of T does not involve θ).

And:

$$T = \frac{\sum X_i}{\sum X_i + \sum Y_j} \sim \text{Beta}(n, m)$$

So the critical values being referenced above may be found via taking critical regions of the Beta distribution when n and m are known.

Q2

Problem 8.28, Casella and Berger (2nd Edition)

Let $f(x|\theta)$ be the logistic location probability density function:

$$f(x|\theta) = \frac{e^{(x-\theta)}}{(1+e^{(x-\theta)})^2}, \quad -\infty < x < \infty, \quad -\infty < \theta < \infty.$$

a)

Show that this family has an MLR.

Answer

Let $\theta_2 > \theta_1$.

We know the likelihood ratio statistic is given by:

$$\Lambda = \frac{f(x|\theta_2)}{f(x|\theta_1)} = e^{\theta_1 - \theta_2} \left[\frac{1+e^{x-\theta_1}}{1+e^{x-\theta_2}} \right]^2$$

The derivative wrt X is of the form:

$$\Lambda' = \frac{e^{x-\theta_1}(1+e^{x-\theta_2}) - e^{x-\theta_2}(1+e^{x-\theta_1})}{(1+e^{x-\theta_2})^2} = \frac{e^{x-\theta_1} - e^{x-\theta_2}}{(1+e^{x-\theta_2})^2} > 0$$

And the inequality holds because of the assumption $\theta_2 > \theta_1$, which is allowed in the full parameter space.

Thus, our likelihood ratio is strictly increasing in x, meaning it is monotonic, i.e. that the family $\{f(x|\theta)\}$ has MLR in x.

b)

Based on one observation X, find the most powerful size α test of

$$H_0 : \theta = 0 \quad \text{versus} \quad H_1 : \theta = 1.$$

For $\alpha = 0.2$, find the size of the Type II error.

Answer

By the Neyman-Pearson Lemma, the MP test rejects H_0 when:

$$\Lambda = \frac{f(x|1)}{f(x|0)} = e^{-1} \left(\frac{1+e^x}{1+e^{x-1}} \right)^2 > k$$

From from part a), since the likelihood ratio is increasing in x, the MP test rejects if $X > k_1$, where k_1 is determined by the size α.

As we know the underlying distributions, let us consider the CDF of the logistic distribution:

$$F(x|\theta) = \frac{e^{x-\theta}}{1+e^{x-\theta}}$$

Under H_0, the size condition is:

$$P(X > k_1 \mid \theta = 0) = 1 - F(k_1|0) = \frac{1}{1+e^{k_1}} = \alpha.$$

Solving for k_1:

$$k_1 = \log\left(\frac{1-\alpha}{\alpha}\right) = \log(\alpha^{-1} - 1)$$

For $\alpha = 0.2$:

$$k_1 = \log(0.2^{-1} - 1) = \log(4) \approx 1.386$$

Under H_1:

$$\beta = P(X \leq k_1 \mid \theta = 1) = F(k_1|1) = \frac{e^{k_1-1}}{1+e^{k_1-1}} \approx \frac{e^{0.386}}{1+e^{0.386}} \approx 0.595$$

So, the MP level test of size $\alpha = 0.2$ rejects when our single observation $X > 1.386$, with a Type II error rate of 0.595.

c)

Show that the test in part b) is UMP size α for testing

$$H_0 : \theta \leq 0 \quad \text{versus} \quad H_1 : \theta > 0.$$

What can be said about UMP tests in general for the logistic location family?

Answer

Via MLR: From part a), the family has MLR in X.

Via Karlin-Rubin Thm. (Knew it would come up again!): Since the MP test for $\theta = 0$ vs $\theta = 1$ rejects for large X and does not depend on the specific parameter value, i.e., $\theta_1 = ...$, it depends solely upon the observed value X, the MP test is also the UMP test for $H_0 : \theta \leq 0$ vs $H_1 : \theta > 0$.

The above results extend to similar distributions within the the logistic location family, i.e., UMP tests for one-sided hypotheses both exist and take the form "Reject H_0 if $X > c$."

Q3

Problem 8.29 a) - b), Casella and Berger (2nd Edition)

Let X be one observation from a Cauchy(θ) distribution.

The Cauchy(θ) density is given by:

$$f(x|\theta) = \frac{1}{\pi} \cdot \frac{1}{1+(x-\theta)^2}, \quad x \in \mathbb{R}, -\infty < \theta < \infty.$$

a)

Show that this family does not have an MLR.

Hint:

Show that the Cauchy(θ) family $\{f(x|\theta) : \theta \in \mathbb{R} = \Theta\}$, based on one observation X, does not have monotone likelihood ratio (MLR) in $t(X) = X$ or $t(X) = -X$. That is, the ratio

$$\frac{f(x|\theta_2)}{f(x|\theta_1)}$$

might not be monotone (either increasing or decreasing) in x.

Answer

Let $\theta_2 > \theta_1$ under the setup of the problem.

The likelihood ratio is of the form:

$$\Lambda = \frac{f(x|\theta_2)}{f(x|\theta_1)} = \frac{1+(x-\theta_1)^2}{1+(x-\theta_2)^2}$$

We know that:

$$\lim_{x\to\pm\infty} \frac{f(x|\theta_2)}{f(x|\theta_1)} = 1$$

So, the ratio achieves a maximum at 1.

As we seek to disprove that the ratio is not monotonic, we need only one example that displays non-monotonicity.

For example, let $\theta_1 = 0$, $\theta_2 = 1$ such that our base assumption that $\theta_2 > \theta_1$ holds.

Then:

$$\Lambda = \frac{f(x|1)}{f(x|0)} = \frac{1+x^2}{1+(x-1)^2}$$

```r
cauchy <- function(x) {
  numerator <- 1 + x^2
  denominator <- 1 + (x-1)^2
  numerator/denominator
}
```

```r
cauchy(0)
cauchy(1)
cauchy(2)
cauchy(3)
cauchy(1000)
```

```r
x_vals <- seq(0, 10, length.out = 500)
y_vals <- cauchy(x_vals)

plot(x_vals, y_vals, type = "l", lwd = 2, col = "black",
     main = "Cauchy Ratio",
     xlab = "x", ylab = expression(Lambda))
```

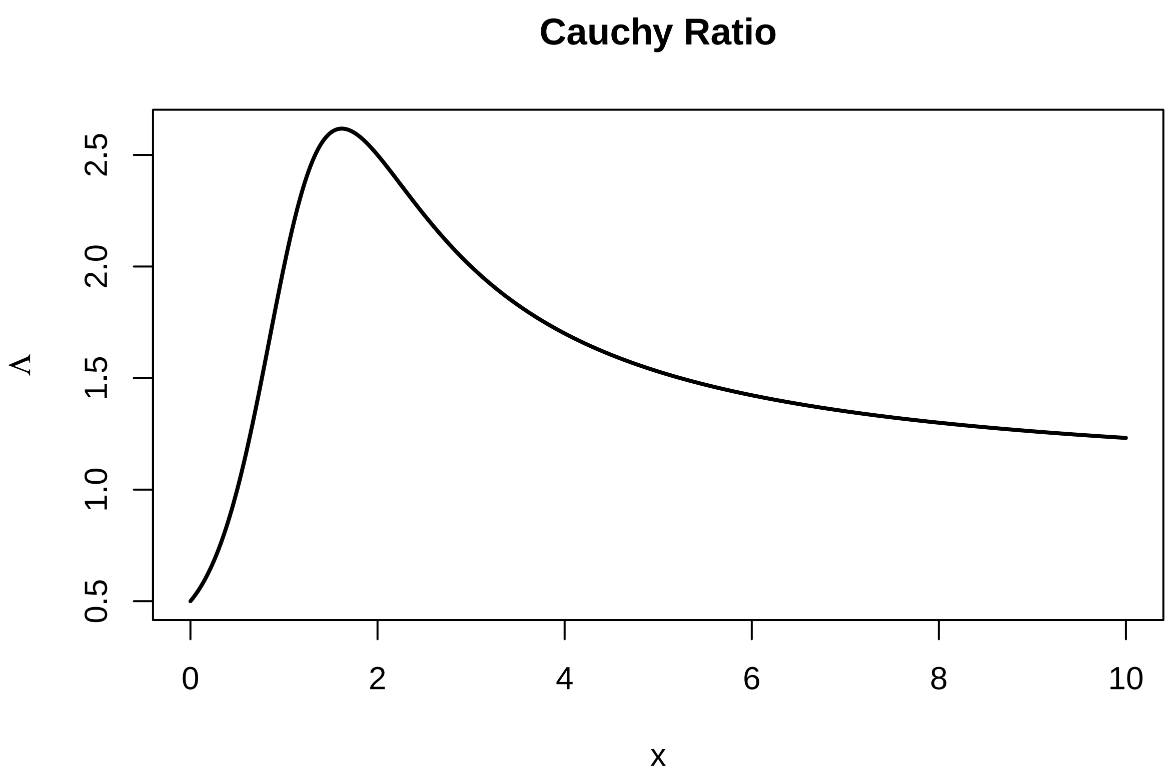

At $x = 0$, $\Lambda = 0.5$. At $x = 1$, $\Lambda = 2$. At $x = 2$, $\Lambda = 2.5$. At $x = 1000$, $\Lambda = 1.002$ (as $x \to \infty$, $\Lambda \to 1$).

The ratio increases from 0 to around 2 and then decreases. So the ratio is not monotonic.

Because the likelihood ratio is not monotonic, then the Cauchy(θ) family lacks MLR in X or $-X$.

b)

Show that the test

$$\phi(x) = \begin{cases} 1 & \text{if } 1 < x < 3 \\ 0 & \text{otherwise} \end{cases}$$

is most powerful of its size for testing

$$H_0 : \theta = 0 \quad \text{versus} \quad H_1 : \theta = 1.$$

Calculate the Type I and Type II error probabilities.

Hint:

Show that the test given is equivalent to rejecting H_0 if

$$f(x|\theta = 1) > 2f(x|\theta = 0)$$

and not rejecting otherwise. Conclude that this must be the most powerful (MP) test for its size. Justify why.

Answer

Consider the test provided in the hint:

$$\phi(x) = \begin{cases} 1 & \text{if } 1 < x < 3, \\ 0 & \text{otherwise} \end{cases}$$

By the Neyman-Pearson Lemma, the MP test rejects H_0 when:

$$\frac{f(x|1)}{f(x|0)} = \frac{1 + x^2}{1 + (x-1)^2} > k$$

We know that the ratio $\frac{f(x|1)}{f(x|0)}$ has critical points at $x = \frac{1 \pm \sqrt{5}}{2}$, because:

$$\Lambda' = \frac{d\Lambda}{dx} = \frac{(2x)(x^2 - 2x + 2) - (1 + x^2)(2x - 2)}{(x^2 - 2x + 2)^2}$$

$$\Lambda' = 0 \to 2x(x^2 - 2x + 2) - (1 + x^2)(2x - 2) = 0 \to x = \frac{1 \pm \sqrt{5}}{2}$$

At any rate, at $x = 1$ and $x = 3$:

$$\frac{f(1|1)}{f(1|0)} = \frac{f(3|1)}{f(3|0)} = 2$$

And the set $\{x : \frac{f(x|1)}{f(x|0)} > 2\} = (1, 3)$ exactly matches the closed form expression of $\phi(x)$.

Since these are one and the same, then $\phi(x)$ is the most powerful test for its size.

Let us then consider the hypotheses we're dealing with.

Under H_0:

$$\alpha = P(1 < X < 3 \mid \theta = 0) = \frac{1}{\pi}\left(\tan^{-1}(3) - \tan^{-1}(1)\right) \approx 0.1476$$

Under H_1:

$$\beta = 1 - P(1 < X < 3 \mid \theta = 1) = 1 - \frac{1}{\pi}\left(\tan^{-1}(2) - \tan^{-1}(0)\right) \approx 0.6476$$

So $\phi(x)$ is MP with $\alpha \approx 0.1476$ (Type I Error Rate) and $\beta \approx 0.6476$ (Type II Error Rate).

Q4

Consider one observation X from the probability density function

$$f(x \mid \theta) = 1 - \theta^2 \left(x - \frac{1}{2} \right), \quad 0 \le x \le 1, \quad 0 \le \theta \le 1.$$

We wish to test:

$$H_0 : \theta = 0 \quad \text{vs.} \quad H_1 : \theta > 0$$

a)

Find the UMP test of size $\alpha = 0.05$ based on X. Carefully justify your answer.

Answer

To find a UMP test of size $\alpha = 0.05$, we first turn to the question of monotonicity, specifically to determine if the family of distribution has MLR in x.

Let $\theta_2 > \theta_1 \ge 0$, as we do.

Then, the likelihood ratio is given by:

$$\Lambda = \frac{f(x \mid \theta_2)}{f(x \mid \theta_1)} = \frac{1 - \theta_2^2(x - \frac{1}{2})}{1 - \theta_1^2(x - \frac{1}{2})}$$

To analyze monotonicity of the ratio, consider first that $x = \frac{1}{2}$ is a critical point. So, we analyze the behavior of the ratio when x is less than or greater than the critical point. Specifically:

When $x < \frac{1}{2}$: both numerator and denominator are increasing in x, and the ratio increases.

When $x > \frac{1}{2}$: both numerator and denominator are decreasing in x, and the ratio decreases.

So unfortunately, Λ is not monotonic across the domain, i.e., the direction of monotonicity changes at $x = \frac{1}{2}$. And consequently the family does not MLR in x (no Karlin-Rubin to make things easier).

However, we can still turn to the Neyman–Pearson Lemma to construct the most powerful test for:

$$H_0 : \theta = 0 \quad \text{vs.} \quad H_1 : \theta = \theta_1 > 0$$

Under H_0:

$$f(x \mid 0) = 1, \quad \text{so } X \sim \text{Uniform}(0, 1)$$

Under H_1:

$$f(x \mid \theta_1) = 1 - \theta_1^2 \left(x - \frac{1}{2} \right)$$

This density under the alternative hypothesis is greater than 1 when $x < \frac{1}{2}$, and less than 1 when $x > \frac{1}{2}$. So $f(x \mid \theta_1)$ has more mass in the tails (0 and 1) than under H_0. So we'd expect the LRT to be much larger (tend to reject H_0) for observations near the tails of the distribution (near 0 or near 1).

So, generally, the most powerful test at level α will reject H_0 for values of x far from $1/2$, i.e. when values are "close" to 0 or "close" to 1.

Under H_0, $X \sim \text{Unif}(0,1)$, we construct a level-$\alpha = 0.05$ test that rejects for extreme values of x.

We define our test function as:

$$\varphi(x) = \begin{cases} 1 & \text{if } x < c_1 \text{ or } x > c_2 \\ 0 & \text{otherwise} \end{cases}$$

Noting that X is a continuous random variable such that our "coin flip" case in the above test is able to be simplified to $\gamma = 0$.

At any rate, our next objective is to further specify the values of c_1 and c_2 such that:

$$P_{\theta=0}(X < c_1) + P_{\theta=0}(X > c_2) = 0.05$$

Since the Uniform(0,1) (distribution under the null) is symmetric, we can divide the error equally to both sides of the density, i.e.,:

$$c_1 = \frac{\alpha}{2} = 0.025, \quad \text{and } c_2 = 1 - \frac{\alpha}{2} = 0.975$$

Giving us the test function of the form:

$$\varphi(x) = \begin{cases} 1 & \text{if } x < 0.025 \text{ or } x > 0.975 \\ 0 & \text{otherwise} \end{cases}$$

This test has size $\alpha = 0.05$, and by the Neyman–Pearson Lemma since the above MP test does not depend on the parameter θ_1, it is also the UMP test of size $\alpha = 0.05$ for testing $H_0 : \theta = 0$ vs. $H_1 : \theta > 0$.

b)

Find the likelihood ratio test statistic $\lambda(X)$ based on X, expressed as a function of X.

Answer

The LRT is:

$$\lambda(X) = \frac{f(X \mid 0)}{\max_{\theta \in [0,1]} f(X \mid \theta)} = \frac{1}{\max_\theta \left[1 - \theta^2 (X - \frac{1}{2})\right]}$$

Again, our critical value is at $\frac{1}{2}$, so we consider the behavior of the LRT at the value of, less than, and greater than $x = \frac{1}{2}$.

For $X \geq \frac{1}{2}$, the maximum occurs at $\theta = 0$, i.e.,:

$$\max_\theta f(X \mid \theta) = 1$$

For $X < \frac{1}{2}$, the maximum occurs at $\theta = 1$:

$$\max_\theta f(X \mid \theta) = 1 + \left(\frac{1}{2} - X\right) = 1.5 - X$$

Incorporating the two cases together, our LRT is of the form:

$$\lambda(X) = \begin{cases} \frac{1}{1.5-X} & \text{if } X < \frac{1}{2} \\ 1 & \text{if } X \geq \frac{1}{2} \end{cases}$$

c)

Find the likelihood ratio test (LRT) of size $\alpha = 0.05$ for the above hypotheses.

Answer

Rejection Region: From part b), $\lambda(X) = 1$ for $X \geq \frac{1}{2}$, and is increasing for $X < \frac{1}{2}$. So to make the test most powerful while maintaining the correct size, we reject for large values of X, where the "large values" are determined by the size condition, which is:

$$P_{\theta=0}(X > k) = 1 - k = 0.05 \quad \Rightarrow \quad k = 0.95$$

Taken together, we reject H_0 when $X > 0.95$. So the test of size $\alpha = 0.05$ is given by:

$$\varphi(x) = \begin{cases} 1 & \text{if } x > 0.95 \\ 0 & \text{otherwise} \end{cases}$$