Actual HW3

Sam Olson

Q1

Using the mixture model (1) and (2) fit the data in Data Set 1 and Data Set 2 separately using maximum likelihood. Produce approximate 95% confidence intervals for the modes of the distributions – these are the value of the parameter ξ in (2). Conduct any additional analysis needed to construct a likelihood ratio test of a reduced model in which the parameters ξ and θ are identical for both data sets against a full model in which those parameters differ between the two heights. Graph the extreme value density or densities for whichever model (full or reduced) you select.

Data model:

$$f(y \mid \mu_i) = \frac{1}{(2\pi)^{1/2}} \exp\left[-\frac{1}{2}(y-\mu_i)^2\right]. \tag{1}$$

Mixing distribution:

$$g(\mu \mid \xi, \theta) = \frac{1}{\theta} \exp\left(-\frac{\mu-\xi}{\theta}\right) \exp\left[-\exp\left(-\frac{\mu-\xi}{\theta}\right)\right], \qquad \theta > 0. \tag{2}$$

A:

Direct Likelihood Evaluation

Table 1: Maximum likelihood estimates and Wald 95% confidence intervals for the mode parameter ξ.

Height	$\hat{\xi}$	Lower 95%	Upper 95%
50m	4.392	3.719	5.066
80m	6.717	5.646	7.787

Profile Likelihood Evaluation

Table 2: Ninety-five percent profile likelihood confidence intervals for the mode parameter ξ.

Height	95% Profile Likelihood CI for ξ
50m	[3.700, 5.075]
80m	[5.633, 7.823]

Likelihood Ratio Test

Table 3: Likelihood ratio test of equality of (ξ, θ) across heights (Full vs. Reduced model).

Hypothesis	LR Statistic	df	p-value
H_0: common (ξ, θ) across heights	18.579	2	9.24e-05

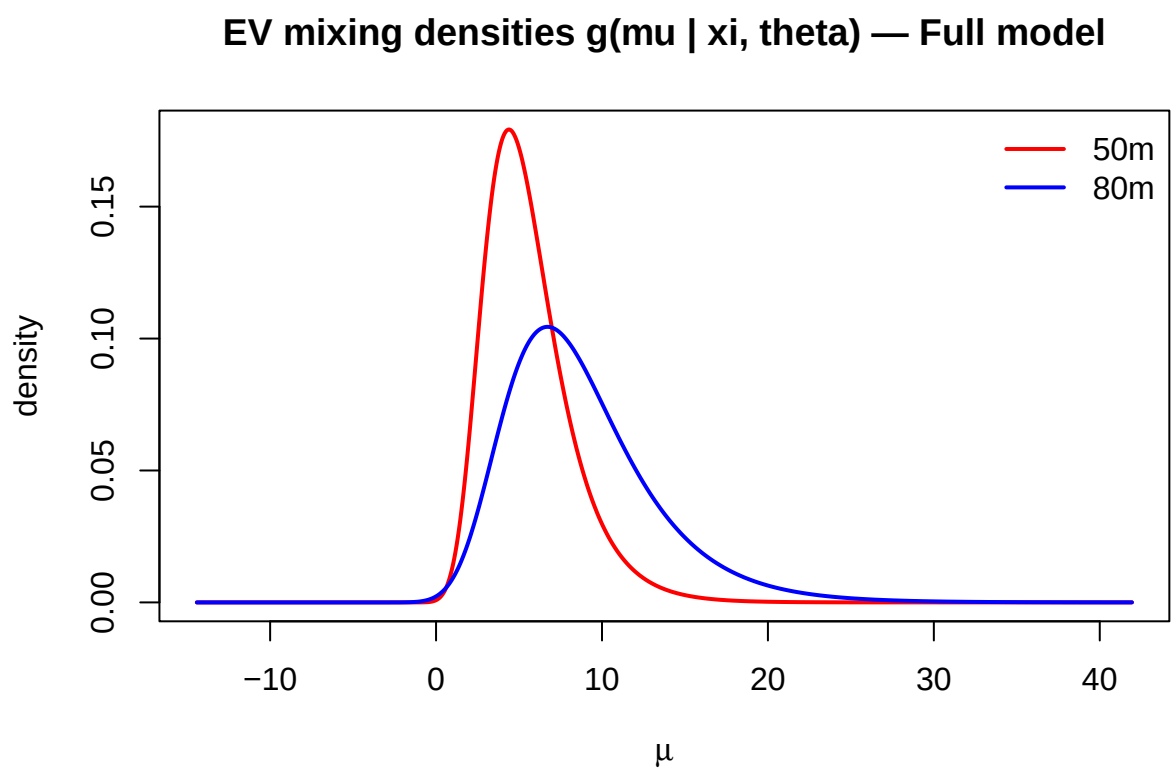

Q2

In (2), $E(\mu) \approx \xi + 0.577\theta$, where 0.577 is Euler's constant (don't confuse with Euler's number). Using this expression, produce maximum likelihood estimate or estimates and 95% confidence interval or confidence intervals for the expected true maximum wind speed from your final model. Note that this will also be the expected observed maximum wind speed.

A:

Table 4: Estimated expected true wind speed $E(\mu)$ and delta-method (Wald) 95% confidence intervals.

Height	$\widehat{E(\mu)}$	SE	Lower 95%	Upper 95%
50m	5.577	0.396	4.802	6.353
80m	8.750	0.651	7.474	10.026

Q3

Find a maximum likelihood estimate or estimates (depending on your final model) of the median daily maximum wind speed or speeds. Compute the associated 95% approximate confidence interval or intervals.

A:

Table 5: Estimated median wind speed and delta-method (Wald) 95% confidence intervals.

Height	$\widehat{\text{Median}}(Y)$	SE	Lower 95%	Upper 95%
50m	5.208	0.367	4.488	5.928
80m	8.049	0.601	6.871	9.227

Q4

Find a maximum likelihood estimate or estimates (depending on your final model) for the probability that true daily maximum wind speed is between 4 and 17 m/s. Compute 95% approximate confidence intervals for the estimate or estimates.

A

Table 6: Estimated probability $Pr(4 \leq \mu \leq 17)$ with 95% Wald confidence intervals.

Height	$\widehat{Pr(4 \leq \mu \leq 17)}$	SE	Lower 95%	Upper 95%
50m	0.700	0.060	0.583	0.817
80m	0.832	0.046	0.743	0.922

Q5

What would you conclude about the question of whether there is or is not a benefit to placing turbine hubs at either 50 or 80 m in height?

A:

In Q1, the likelihood ratio test provided strong evidence that the wind speed distributions at 50m and 80m differ ($p \approx 9.24 \times 10^{-5}$), leading us to adopt the full model with separate parameter estimates for each height. The fitted extreme value mixing distributions therefore represent statistically distinct wind regimes.

Under the selected model, all key summaries indicate systematically higher wind speeds at 80m. The estimated expected true daily maximum wind speed is approximately 3.2 m/s higher at 80m (8.75 vs. 5.58 m/s), and the estimated median differs by a similar margin (8.05 vs. 5.21 m/s). Moreover, the estimated probability that true daily maximum wind speed falls within the operationally relevant range of 4–17 m/s is substantially larger at 80m (0.832 vs. 0.700), with confidence intervals suggesting that this difference is unlikely to be explained by sampling variability alone.

Taken together, these results indicate that, under the fitted mixture model and for wind conditions comparable to those observed, a turbine hub height of 80m is associated with more favorable wind speed characteristics for energy production. This conclusion remains conditional on the correctness of the model assumptions and on future wind patterns being similar to those represented in the data.

Q6

As an exercise in learning about statistical procedures, your friend would like to see a Monte Carlo approximation to the expected true wind speed for the two heights separately. Produce such approximations along with 95% Monte Carlo intervals. Your friend understands intuitively that these approximations and the maximum likelihood estimates of Question 2 should have the same or nearly the same values, but marvels at the fact that the intervals associated with these approximations are so much more narrow than those computed using the asymptotic theory of likelihood estimation in Question 2. Explain why this is the case and what it means.

A:

Monte Carlo

Table 7: Monte Carlo approximation of $E(\mu)$ using M simulated draws, with 95% Monte Carlo intervals.

Height	M	$\widehat{E(\mu)}_{\text{MC}}$	Lower 95%	Upper 95%
50m	20000	5.583	5.546	5.620
80m	20000	8.706	8.644	8.768

Comparison - Concise

The Monte Carlo (MC) intervals are much narrower because they measure only simulation (numerical) error conditional on the estimated parameters $(\hat{\xi}, \hat{\theta})$, whereas the asymptotic confidence intervals in Question 2 account for statistical uncertainty in the joint estimation of (ξ, θ) from finite data.

Comparison - More Details

For each height, we simulated

$$\mu^{(1)}, \ldots, \mu^{(M)} \sim \text{EV}(\hat{\xi}, \hat{\theta})$$

and estimated the expected true wind speed by

$$\widehat{E}_{MC}(\mu) = \bar{\mu}_M = \frac{1}{M} \sum_{m=1}^{M} \mu^{(m)}.$$

By the Central Limit Theorem for simulation,

$$\bar{\mu}_M \approx \mathcal{N}\left(E(\mu), \frac{\text{Var}(\mu)}{M}\right),$$

so an approximate 95% Monte Carlo interval is

$$\bar{\mu}_M \pm 1.96 \frac{s_M}{\sqrt{M}}.$$

The crucial point is that this interval treats $(\hat{\xi}, \hat{\theta})$ as fixed and quantifies only Monte Carlo approximation error, which decreases at rate $1/\sqrt{M}$ and can be made arbitrarily small by increasing M.

In contrast, the confidence interval in Question 2 is based on asymptotic likelihood theory for

$$E(\mu) = \xi + \gamma\theta, \qquad \gamma \approx 0.577,$$

and reflects uncertainty in estimating (ξ, θ) from $n = 50$ observations. That parameter uncertainty propagates into uncertainty about $E(\mu)$.

Thus, the Monte Carlo intervals are narrow because they measure numerical precision, not inferential uncertainty. The wider intervals from Question 2 are the appropriate measures of statistical uncertainty about the expected wind speed. To obtain a simulation-based interval comparable to Question 2, one would need to incorporate parameter uncertainty (e.g., via a parametric bootstrap), rather than conditioning on fixed MLEs.

Appendix

```r
library(tidyverse)
library(numDeriv)

# External optimization helpers (Newton-Raphson, etc.)
source("C:/Users/samue/OneDrive/Desktop/Iowa_State_PS/STAT_6010/PS3/newtraph.txt")
source("C:/Users/samue/OneDrive/Desktop/Iowa_State_PS/STAT_6010/PS3/bisect.txt")

# Read data
m50 <- read.table("C:/Users/samue/OneDrive/Desktop/Iowa_State_PS/STAT_6010/PS3/winddat1.txt",
                  header = TRUE)
m80 <- read.table("C:/Users/samue/OneDrive/Desktop/Iowa_State_PS/STAT_6010/PS3/winddat2.txt",
                  header = TRUE)

y50 <- as.numeric(m50$y)
y80 <- as.numeric(m80$y)

# Euler-Mascheroni constant
gamma <- 0.5772156649015329

# ------------------------------------------------------------
# Extreme value (right-skew Gumbel) density for mu
# ------------------------------------------------------------
g_mu <- function(mu, xi, theta) {
  z <- (mu - xi) / theta
  (1 / theta) * exp(-z - exp(-z))
}

# ------------------------------------------------------------
# Stable marginal density m(y | xi, theta)
#   Uses reparameterization z = (mu - xi)/theta
#   m(y) = int phi(y | xi + theta z, 1) * exp(-z - exp(-z)) dz
# ------------------------------------------------------------
m_y_stable <- function(y, xi, theta, zL = -10, zU = 30) {
  integrand <- function(z) {
    log_int <- dnorm(y, mean = xi + theta * z, sd = 1, log = TRUE) - z - exp(-z)
    exp(log_int)
  }
  val <- integrate(integrand, lower = zL, upper = zU,
                   rel.tol = 1e-8, abs.tol = 0, subdivisions = 5000)$value
  pmax(val, .Machine$double.xmin)
}

# Vectorized log m(y)
log_m_vec_stable <- function(y, xi, theta, zL = -10, zU = 30) {
  vapply(y, function(yy) log(m_y_stable(yy, xi, theta, zL, zU)), numeric(1))
}

# ------------------------------------------------------------
# Log-likelihood for one dataset
# Parameterization: par = (xi, logtheta)
# ------------------------------------------------------------
```

```r
logLik_one <- function(par, y, zL = -10, zU = 30) {
  xi <- par[1]
  theta <- exp(par[2])
  sum(log_m_vec_stable(y, xi, theta, zL = zL, zU = zU))
}

# Negative log-likelihood (for optimization routines)
nll_one <- function(par, y, zL = -10, zU = 30) {
  -logLik_one(par, y = y, zL = zL, zU = zU)
}

# ------------------------------------------------------------
# Moment-based starting values
#   Var(Y) = (pi^2/6) theta^2 + 1
#   E(Y)   = xi + gamma * theta
# ------------------------------------------------------------
ev_start <- function(y) {
  v <- var(y)
  theta0 <- sqrt(pmax((v - 1) * 6 / pi^2, 1e-6))
  xi0 <- mean(y) - gamma * theta0
  c(xi0, log(theta0))
}

# ------------------------------------------------------------
# newtraph() interface: returns (logLik, gradient, Hessian)
# ------------------------------------------------------------
make_ders_one <- function(zL = -10, zU = 30) {
  function(par, dat) {
    f <- function(pp) logLik_one(pp, y = dat, zL = zL, zU = zU)
    ll <- f(par)
    g <- grad(f, par)
    H <- hessian(f, par)
    list(ll, g, H)
  }
}

# ------------------------------------------------------------
# Fit one dataset using Newton-Raphson
# Returns MLEs, Wald CI for xi, and covariance matrix
# ------------------------------------------------------------
fit_one_newtraph <- function(y, zL = -10, zU = 30) {
  start <- ev_start(y)
  ders <- make_ders_one(zL = zL, zU = zU)
  out <- newtraph(ders, dat = y, x0 = start,
                  quiet = TRUE, print_final = FALSE,
                  step_scale = 0.5, max_halving = 60)
  par_hat <- as.numeric(out[[1]])
  xi_hat <- par_hat[1]
  theta_hat <- exp(par_hat[2])
  fInf <- out[[3]]               # (-H)^(-1) at solution
  se_xi <- sqrt(fInf[1, 1])
  ci_xi_wald <- c(xi_hat - 1.96 * se_xi,
                  xi_hat + 1.96 * se_xi)
```

```r
  list(xi_hat = xi_hat,
       theta_hat = theta_hat,
       logLik = out[[2]],
       fInf = fInf,
       se_xi = se_xi,
       ci_xi_wald = ci_xi_wald)
}

fit50 <- fit_one_newtraph(y50, zL = -100, zU = 100)
fit80 <- fit_one_newtraph(y80, zL = -100, zU = 100)

# ------------------------------------------------------------
# Profile likelihood CI for xi
# Solves: 2(ell_hat - ell_prof(xi)) = chi^2_{1,0.95}
# ------------------------------------------------------------
profile_logLik_xi <- function(xi_fixed, y, logtheta_center,
                              zL = -10, zU = 30) {
  obj <- function(logtheta)
    nll_one(c(xi_fixed, logtheta), y = y, zL = zL, zU = zU)
  opt <- optimize(obj,
                  interval = c(logtheta_center - log(20),
                               logtheta_center + log(20)))
  list(logLik = -opt$objective)
}

ci_xi_profile <- function(y, xi_hat, theta_hat, logLik_hat,
                          zL = -10, zU = 30, alpha = 0.05) {
  cutoff <- qchisq(1 - alpha, df = 1)
  target <- function(xi) {
    ell_prof <- profile_logLik_xi(
      xi_fixed = xi,
      y = y,
      logtheta_center = log(theta_hat),
      zL = zL, zU = zU
    )$logLik
    2 * (logLik_hat - ell_prof) - cutoff
  }
  step0 <- max(0.25, 0.5 * theta_hat)
  lower <- uniroot(target,
                   interval = c(xi_hat - 5 * step0, xi_hat))$root
  upper <- uniroot(target,
                   interval = c(xi_hat, xi_hat + 5 * step0))$root
  c(lower, upper)
}

ci50_prof <- ci_xi_profile(y50, fit50$xi_hat,
                           fit50$theta_hat, fit50$logLik,
                           zL = -100, zU = 100)

ci80_prof <- ci_xi_profile(y80, fit80$xi_hat,
                           fit80$theta_hat, fit80$logLik,
                           zL = -100, zU = 100)
```

```r
# ------------------------------------------------------------
# Likelihood ratio test: full vs reduced model
# Full: separate (xi, theta) by height
# Reduced: common (xi, theta)
# ------------------------------------------------------------
logLik_full <- function(par, y50, y80,
                        zL = -100, zU = 100) {
  xi1 <- par[1]; th1 <- exp(par[2])
  xi2 <- par[3]; th2 <- exp(par[4])
  sum(log_m_vec_stable(y50, xi1, th1, zL, zU)) +
    sum(log_m_vec_stable(y80, xi2, th2, zL, zU))
}

logLik_reduced <- function(par, y50, y80,
                           zL = -100, zU = 100) {
  xi <- par[1]; th <- exp(par[2])
  sum(log_m_vec_stable(y50, xi, th, zL, zU)) +
    sum(log_m_vec_stable(y80, xi, th, zL, zU))
}

full_opt <- optim(c(fit50$xi_hat, log(fit50$theta_hat),
                    fit80$xi_hat, log(fit80$theta_hat)),
                  fn = function(p) -logLik_full(p, y50, y80),
                  method = "BFGS")

red_opt <- optim(ev_start(c(y50, y80)),
                 fn = function(p) -logLik_reduced(p, y50, y80),
                 method = "BFGS")

LR <- 2 * (-full_opt$value - (-red_opt$value))
pval <- 1 - pchisq(LR, df = 2)
use_full <- (pval < 0.05)

# Plot EV mixing densities for selected model
if (use_full) {
  xi50_hat <- full_opt$par[1]
  th50_hat <- exp(full_opt$par[2])
  xi80_hat <- full_opt$par[3]
  th80_hat <- exp(full_opt$par[4])
  mu_grid <- seq(min(xi50_hat - 6 * th50_hat,
                     xi80_hat - 6 * th80_hat),
                 max(xi50_hat + 10 * th50_hat,
                     xi80_hat + 10 * th80_hat),
                 length.out = 800)
  plot(mu_grid, g_mu(mu_grid, xi50_hat, th50_hat),
       type = "l", col = "red", lwd = 2,
       xlab = expression(mu), ylab = "density")
  lines(mu_grid, g_mu(mu_grid, xi80_hat, th80_hat),
        col = "blue", lwd = 2)
} else {
  xi_hat <- red_opt$par[1]
  th_hat <- exp(red_opt$par[2])
  mu_grid <- seq(xi_hat - 6 * th_hat,
```

```r
                 xi_hat + 10 * th_hat,
                 length.out = 800)
  plot(mu_grid, g_mu(mu_grid, xi_hat, th_hat),
       type = "l", lwd = 2,
       xlab = expression(mu), ylab = "density")
}

# ------------------------------------------------------------
# Q2: E(mu) = xi + gamma * theta
# Delta-method CI
# ------------------------------------------------------------
wald_ci_mean_mu <- function(fit) {
  xi_hat <- fit$xi_hat
  th_hat <- fit$theta_hat
  V <- fit$fInf
  eta_hat <- xi_hat + gamma * th_hat
  grad_vec <- c(1, gamma * th_hat)
  se <- sqrt(drop(t(grad_vec) %*% V %*% grad_vec))
  ci <- eta_hat + c(-1, 1) * qnorm(0.975) * se
  list(eta_hat = eta_hat, se = se, ci95 = ci)
}

ci50_eta_wald <- wald_ci_mean_mu(fit50)
ci80_eta_wald <- wald_ci_mean_mu(fit80)

# ------------------------------------------------------------
# Q3: Median of observed Y under mixture
# Solve F_Y(y) = 0.5
# ------------------------------------------------------------
FY_stable <- function(y, xi, theta,
                      zL = -100, zU = 100) {
  integrand <- function(z)
    pnorm(y - (xi + theta * z)) *
    exp(-z - exp(-z))
  val <- integrate(integrand, lower = zL, upper = zU)$value
  pmin(pmax(val, 0), 1)
}

median_y <- function(xi, theta) {
  f <- function(y) FY_stable(y, xi, theta) - 0.5
  uniroot(f,
          interval = c(xi - 10 * theta,
                       xi + 10 * theta))$root
}

wald_ci_median_y <- function(fit) {
  xi_hat <- fit$xi_hat
  th_hat <- fit$theta_hat
  V <- fit$fInf
  h <- function(par) {
    xi <- par[1]
    th <- exp(par[2])
    median_y(xi, th)
```

```r
  }
  par_hat <- c(xi_hat, log(th_hat))
  m_hat <- h(par_hat)
  g <- grad(h, par_hat)
  se <- sqrt(drop(t(g) %*% V %*% g))
  ci <- m_hat + c(-1, 1) * qnorm(0.975) * se
  list(median_hat = m_hat, se = se, ci95 = ci)
}

ci50_med <- wald_ci_median_y(fit50)
ci80_med <- wald_ci_median_y(fit80)

# ------------------------------------------------------------
# Q4: P(4 <= mu <= 17) under EV mixing distribution
# ------------------------------------------------------------
G_mu <- function(u, xi, theta)
  exp(-exp(-(u - xi) / theta))

wald_ci_p_between_mu <- function(fit, a = 4, b = 17) {
  xi_hat <- fit$xi_hat
  th_hat <- fit$theta_hat
  V <- fit$fInf
  h <- function(par) {
    xi <- par[1]
    th <- exp(par[2])
    G_mu(b, xi, th) - G_mu(a, xi, th)
  }
  par_hat <- c(xi_hat, log(th_hat))
  p_hat <- h(par_hat)
  g <- grad(h, par_hat)
  se <- sqrt(drop(t(g) %*% V %*% g))
  ci <- p_hat + c(-1, 1) * qnorm(0.975) * se
  list(p_hat = p_hat, se = se, ci95 = ci)
}

ci50_p <- wald_ci_p_between_mu(fit50)
ci80_p <- wald_ci_p_between_mu(fit80)

# ------------------------------------------------------------
# Q6: Monte Carlo approximation of E(mu)
# ------------------------------------------------------------
set.seed(43)

rEV <- function(n, xi, theta) {
  u <- runif(n)
  xi - theta * log(-log(u))
}

mc_E_mu <- function(fit, M = 20000) {
  xi_hat <- fit$xi_hat
  th_hat <- fit$theta_hat
  mu_draws <- rEV(M, xi_hat, th_hat)
  mc_mean <- mean(mu_draws)
```

```r
  mc_se <- sd(mu_draws) / sqrt(M)
  mc_ci <- mc_mean + c(-1, 1) * qnorm(0.975) * mc_se
  plug_in <- xi_hat + gamma * th_hat
  list(mc_mean = mc_mean,
       mc_ci95 = mc_ci,
       plug_in = plug_in)
}

mc50 <- mc_E_mu(fit50)
mc80 <- mc_E_mu(fit80)
```